Documentation for the Combine (focused) crawling system

Anders Ardö, Koraljka Golub

October 13, 2008

Contents

Part I
Overview

1 Introduction

The Combine system is an open, free, and highly configurable system for focused crawling of Internet resources. It aims at providing a robust and efficient tool for creating topic-specific moderate sized databases (up to a few million records). Crawling speed is around 200 URLs per minute and a complete structured record takes up an average of 25 kilobytes disk-space.

Figure 1: Overview of the Combine focused crawler.

Main features include:

- part of the SearchEngine-in-a-Box[1] system

- extensive configuration possibilities

- integrated topic filter (automated topic classifier) for focused crawling mode

- possibility to use any topic filter (if provided as a Perl Plug-In module[2]) in focused crawling mode

- crawl limitations based on regular expression on URLs - both include and exclude rules (URL focus filter)

- character set detection/normalization

- language detection

- HTML cleaning

---

[1] http://combine.it.lth.se/SearchEngineBox/
[2] http://combine.it.lth.se/PlugIns/

- metadata extraction

- duplicate detection

- HTML parsing to provide structured records for each crawled page

- support for many document formats (text, HTML, PDF, PostScript, MSWord, MSPowerPoint, MSExcel, RTF, TeX, images)

- SQL database for data storage and administration

Naturally it obeys the Robots Exclusion Protocol[3] and behaves nice to Web-servers. Besides focused crawls (generating topic-specific databases), Combine supports configurable rules on what's crawled based on regular expressions on URLs (URL focus filter). The crawler is designed to run continuously in order to keep crawled databases as up to date as possible. It can be stopped and restarted any time without loosing any status or information.

The operation of Combine (overview in Figure **??**) as a focused crawler is based on a combination of a general Web crawler and an automated subject classifier. The topic focus is provided by a focus filter using a topic definition implemented as a thesaurus, where each term is connected to a topic class.

Crawled data are stored as a structured records in a local relational database.

Section **??** outlines how to download, install and test the Combine system and includes use scenarios – useful in order to get a jump start at using the system.

Section **??** discusses configuration structure and highlights a few important configuration variables.

Section **??** describes policies and methods used by the crawler.

Evaluation and performance are treated in sections **??** and **??**.

The system has a number of components (see section **??**), the main ones visible to the user being `combineCtrl` which is used to start and stop crawling and view crawler status, and `combineExport` that extracts crawled data from the internal database and exports them as XML records.

Further details (lots and lots of them) can be found in part **??** 'Gory details' and in Appendix **??**.

2   Open source distribution, installation

The focused crawler has been restructured and packaged as a Debian package in order to ease distribution and installation. The package contains dependency information to make sure that all software that is needed to run the crawler is installed at the same time. In connection with this we have also packaged a number of necessary Perl-modules as Debian packages.

All software and packages are available from a number of places:

- the Combine focused crawler Web-site[4]

- the Comprehensive Perl Archive Network - CPAN[5]

---

[3]http://www.robotstxt.org/wc/exclusion.html
[4]http://combine.it.lth.se/
[5]http://search.cpan.org/ aardo/Combine/

- SourceForge project "Combine focused crawler"[6]

In addition to the distribution sites there is a public discussion list at SourceForge[7].

2.1   Installation

This distribution is developed and tested on Linux systems. It is implemented entirely in Perl and uses the MySQL[8] database system, both of which are supported on many other operating systems. Porting to other UNIX dialects should be easy.

The system is distributed either as source or as a Debian package.

2.1.1   Installation from source for the impatient

Unless you are on a system supporting Debian packages (in which case look at Automated installation (section **??**)), you should download and unpack the source. The following command sequence will then install Combine:

```
perl Makefile.PL
make
make test
make install
mkdir /etc/combine
cp conf/* /etc/combine/
mkdir /var/run/combine
```

Test that it all works (run as root)
```
./doc/InstallationTest.pl
```

2.1.2   Porting to not supported operating systems - dependencies

In order to port the system to another platform, you have to verify the availability, for this platform, of the two main systems:

- Perl[9]

- MySQL version ≥ 4.1[10]

If they are supported you stand a good chance to port the system.

Furthermore, the external Perl modules (listed in **??**) should be verified to work on the new platform.

Perl modules are most easily installed using the Perl CPAN automated system
(`perl -MCPAN -e shell`).
Optionally the following external programs will be used if they are installed on your system:

- antiword (parsing MSWord files)

- detex (parsing TeX files)

---

[6] http://sourceforge.net/projects/focused-crawler
[7] http://lists.sourceforge.net/lists/listinfo/focused-crawler-general
[8] http://www.mysql.com/
[9] http://www.cpan.org/ports/index.html
[10] http://dev.mysql.com/downloads/

- pdftohtml (parsing PDF files)

- pstotext (parsing PS and PDF files, needs ghostview)

- xlhtml (parsing MSExcel files)

- ppthtml (parsing MSPowerPoint files)

- unrtf (parsing RTF files)

- tth (parsing TeX files)

- untex (parsing TeX files)

2.1.3 Automated Debian/Ubuntu installation

- Add the following lines to your /etc/apt/sources.list:
  `deb http://combine.it.lth.se/ debian/`

- Give the commands:
  `apt-get update`
  `apt-get install combine`

This also installs all dependencies such as MySQL and a lot of necessary Perl modules.

2.1.4 Manual installation

Download the latest distribution[11].

Install all software that Combine depends on (see above).

Unpack the archive with `tar zxf`

This will create a directory named `combine-XX` with a number of subdirectories including `bin, Combine, doc, and conf`.

'`bin`' contains the executable programs.

'`Combine`' contains needed Perl modules. They should be copied to where Perl will find them, typically `/usr/share/perl5/Combine/`.

'`conf`' contains the default configuration files. Combine looks for them in `/etc/combine/` so they need to be copied there.

'`doc`' contains documentation.

The following command sequence will install Combine:

```
perl Makefile.PL
make
make test
make install
mkdir /etc/combine
cp conf/* /etc/combine/
mkdir /var/run/combine
```

---

[11]http://combine.it.lth.se/#downloads

2.1.5 Out-of-the-box installation test

A simple way to test your newly installed Combine system is to crawl just one Web-page and export it as an XML-document. This will exercise much of the code and guarantee that basic focused crawling will work.

- Initialize a crawl-job named aatest. This will create and populate the job-specific configuration directory and create the MySQL database that will hold the records:

```
sudo combineINIT --jobname aatest --topic /etc/combine/Topic_carnivor.txt
```

- Harvest the test URL by:

```
combine --jobname aatest
        --harvest http://combine.it.lth.se/CombineTests/InstallationTest.html
```

- Export a structured Dublin Core record by:

```
combineExport --jobname aatest --profile dc
```

- and verify that the output, except for dates and order, looks like:

```
<?xml version="1.0" encoding="UTF-8"?>
<documentCollection version="1.1" xmlns:dc="http://purl.org/dc/elements/1.1/">
<metadata xmlns:dc="http://purl.org/dc/elements/1.1/">
<dc:format>text/html</dc:format>
<dc:format>text/html; charset=iso-8859-1</dc:format>
<dc:subject>Carnivorous plants</dc:subject>
<dc:subject>Drosera</dc:subject>
<dc:subject>Nepenthes</dc:subject>
<dc:title transl="yes">Installation test for Combine</dc:title>
<dc:description></dc:description>
<dc:date>2006-05-19 9:57:03</dc:date>
<dc:identifier>http://combine.it.lth.se/CombineTests/InstallationTest.html</dc:identifier>
<dc:language>en</dc:language>
</metadata>
```

Or run – as root – the script ./doc/InstallationTest.pl (see ?? in the Appendix) which essentially does the same thing.

2.2 Getting started

A simple example work-flow for a trivial crawl job name 'aatest' might look like:

1. Initialize database and configuration (needs root privileges)
   `sudo combineINIT --jobname aatest`

2. Load some seed URLs like (you can repeat this command with different URLs as many times as you wish)
   `echo 'http://combine.it.lth.se/' | combineCtrl load --jobname aatest`

7

3. Start 2 harvesting processes
   `combineCtrl start --jobname aatest --harvesters 2`

4. Let it run for some time. Status and progress can be checked using the program
   '`combineCtrl --jobname aatest`' with various parameters.

5. When satisfied kill the crawlers
   `combineCtrl kill --jobname aatest`

6. Export data records in the ALVIS XML format
   `combineExport --jobname aatest --profile alvis`

7. If you want to schedule a recheck for all the crawled pages stored in the database do
   `combineCtrl reharvest --jobname aatest`

8. Go back to **??** for continuous operation.

Once a job is initialized it is controlled using `combineCtrl`. Crawled data is exported using `combineExport`.

2.3   Online documentation

The latest, updated, detailed documentation is always available online[12].

2.4   Use scenarios

2.4.1   General crawling without restrictions

Use the same procedure as in section **??**. This way of crawling is not recommended for the Combine system since it will generate really huge databases without any focus.

2.4.2   Focused crawling – domain restrictions

Create a focused database with all pages from a Web-site. In this use scenario we will crawl the Combine site and the ALVIS site. The database is to be continuously updated, i.e. all pages have to be regularly tested for changes, deleted pages should be removed from the database, and newly created pages added.

1. Initialize database and configuration
   `sudo combineINIT --jobname focustest`

2. Edit the configuration to provide the desired focus
   Change the `<allow>` part in `/etc/combine/focustest/combine.cfg` from

   ```
   #use either URL or HOST: (obs ':') to match regular expressions to either the
   #full URL or the HOST part of a URL.
   <allow>
   #Allow crawl of URLs or hostnames that matches these regular expressions
   HOST: .*$
   </allow>
   ```

---

[12]http://combine.it.lth.se/documentation/

to

```
#use either URL or HOST: (obs ':') to match regular expressions to either the
#full URL or the HOST part of a URL.
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: www\.alvis\.info$
HOST: combine\.it\.lth\.se$
</allow>
```

The escaping of '.' by writing '\.' is necessary since the patterns actually are Perl regular expressions. Similarly the ending '$' indicates that the host string should end here, so for example a Web server on `www.alvis.info.com` (if such exists) will not be crawled.

3. Load seed URLs
   ```
   echo 'http://combine.it.lth.se/' | combineCtrl load --jobname focustest
   echo 'http://www.alvis.info/' | combineCtrl load --jobname focustest
   ```

4. Start 1 harvesting process
   ```
   combineCtrl start --jobname focustest
   ```

5. Daily export all data records in the ALVIS XML format
   ```
   combineExport --jobname focustest --profile alvis
   ```
   and schedule all pages for re-harvesting
   ```
   combineCtrl reharvest --jobname focustest
   ```

2.4.3 Focused crawling – topic specific

Create and maintain a topic specific crawled database for the topic 'Carnivorous plants'.

1. Create a topic definition (see section **??**) in a local file named `cpTopic.txt`. (Can be done by copying `/etc/combine/Topic_carnivor.txt` since it happens to be just that.)

2. Create a file named `cpSeedURLs.txt` with seed URLs for this topic, containing the URLs:

   ```
   http://www.sarracenia.com/faq.html
   http://dmoz.org/Home/Gardening/Plants/Carnivorous_Plants/
   http://www.omnisterra.com/bot/cp_home.cgi
   http://www.vcps.au.com/
   http://www.murevarn.se/links.html
   ```

3. Initialization
   ```
   sudo combineINIT --jobname cptest --topic cpTopic.txt
   ```
   This enables topic checking and focused crawl mode by setting configuration variable `doCheckRecord = 1` and copying a topic definition file (`cpTopic.txt`) to `/etc/combine/cptest/topicdefinition.txt`.

9

4. Load seed URLs
   ```
   combineCtrl load --jobname cptest < cpSeedURLs.txt
   ```

5. Start 3 harvesting process
   ```
   combineCtrl start --jobname cptest --harvesters 3
   ```

6. Regularly export all data records in the ALVIS XML format
   ```
   combineExport --jobname cptest --profile alvis
   ```

Running this crawler for an extended period will result in more than 200 000 records.

2.4.4 Focused crawling in an Alvis system

Use the same procedure as in section **??** (Focused crawling – topic specific) except for the last point. Exporting should be done incrementally into an Alvis pipeline (in this example listening at port 3333 on the machine nlp.alvis.info):
```
combineExport --jobname cptest --pipehost nlp.alvis.info --pipeport 3333 --incremental
```

2.4.5 Crawl one entire site and it's outlinks

This scenario requires the crawler to:

- crawl an entire target site

- crawl all the outlinks from the site

- crawl no other site or URL apart from external URLs mentioned on the one target site

I.e. all of `http://my.targetsite.com/*`, plus any other URL that is linked to from a page in `http://my.targetsite.com/*`.

1. Configure Combine to crawl this one site only. Change the `<allow>` part in `/etc/combine/XXX/combine.cfg` to

   ```
   #use either URL or HOST: (obs ':') to match regular expressions to either the
   #full URL or the HOST part of a URL.
   <allow>
   #Allow crawl of URLs or hostnames that matches these regular expressions
   HOST: my\.targetsite\.com$
   </allow>
   ```

2. Crawl until you have the entire site (if it's a big site you might want to do the changes suggested in FAQ no **??**).

3. Stop crawling.

4. Change configuration `<allow>` back to allow crawling of any domain (which is the default).

```
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: .*$
</allow>
```

5. Schedule all links in the database for crawling, something like (change XXX to your jobname)
```
echo 'select urlstr from urls;' | mysql -u combine XXX
             | combineCtrl load --jobname XXX
```

6. Change configuration to disable automatic recycling of links:
```
#Enable(1)/disable(0) automatic recycling of new links
AutoRecycleLinks = 0
```
and maybe (depending or your other requirements) change:
```
#User agent handles redirects (1) or treat redirects as new links (0)
UserAgentFollowRedirects = 0
```

7. Start crawling and run until no more in queue.

3   Configuration

Configuration files use a simple format consisting of either name/value pairs or complex variables in sections. Name/value pairs are encoded as single lines formated like 'name = value'. Complex variables are encoded as multiple lines in named sections delimited as in XML, using '<name> ... </name>'. Sections may be nested for related configuration variables. Empty lines and lines starting with '#' (comments) are ignored.

The most important configuration variables are the complex variables <url><allow> (allows certain URLs to be harvested) and <url><exclude> (excludes certain URLs from harvesting) which are used to limit your crawl to just a section of the WWW, based on the URL. Loading URLs to be crawled into the system checks each URL first against the Perl regular expressions of <url><allow> and if it matches goes on to match it against <url><exclude> where it's discarded if it matches, otherwise it's scheduled for crawling. (See section **??** 'URL filtering').

3.1   Configuration files

All configuration files are stored in the /etc/combine/ directory tree. All configuration variables have reasonable defaults (section **??**).

3.1.1   Templates

The values in

job_default.cfg contains job specific defaults. It is copied to a subdirectory named after the job by combineINIT.

SQLstruct.sql contains structure of the internal SQL database used both for administration and for holding data records. Details in section **??**.

Topic_* contains various contributed topic definitions.

3.1.2 Global configuration files

Files used for global parameters for all crawler jobs.

default.cfg is the global defaults. It is loaded first. Consult section **??** and appendix **??** for details. Values can be overridden from the job-specific configuration file `combine.cfg`.

tidy.cfg configuration for Tidy cleaning of HTML code.

3.1.3 Job specific configuration files

The program `combineINIT` creates a job specific sub-directory in `/etc/combine` and populates it with some files including `combine.cfg` initialized with a copy of `job_default.cfg`. You should always change the value of the variable `Operator-Email` in this file and set it to something reasonable. It is used by Combine to identify you to the crawled Web-servers.

The job-name have to be given to all programs when started using the `--jobname` switch.

combine.cfg the job specific configuration. It is loaded second and overrides the global defaults. Consult section **??** and appendix **??** for details.

topicdefinition.txt contains the topic definition for focused crawl if the `--topic` switch is given to `combineINIT`. The format of this file is described in section **??**.

stopwords.txt a file with words to be excluded from the automatic topic classification processing. One word per line. Can be empty (default) but must be present.

config_exclude contains more exclude patterns. Optional, automatically included by `combine.cfg`. Updated by `combineUtil`.

config_serveralias contains patterns for resolving Web server aliases. Optional, automatically included by `combine.cfg`. Updated by `combineUtil`.

sitesOK.txt optionally used by the built-in automated classification algorithms (section **??**) to bypass the topic filter for certain sites.

3.1.4 Details and default values

Further details are found in section **??** 'Configuration variables' which lists all variables and their default values.

4 Crawler internal operation

The system is designed for continuous operation. The harvester processes a URL in several steps as detailed in Figure **??**. As a start-up initialization the frontier has to be seeded with some relevant URLs. All URLs are normalized before they are entered in the database. Data can be exported in various formats including the ALVIS XML document format[13] and Dublin Core[14] records.

The steps taken during crawling (numbers refer to Figure **??**):

---

[13]http://www.alvis.info/alvis/architecture
[14]http://dublincore.org/

Figure 2: Architecture for the Combine focused crawler.

1. The next URL is fetched from the scheduler.

2. Combine obeys the Robots Exclusion Protocol[15]. Rules are cached locally.

3. The page is retrieved using a GET, GET-IF-MODIFIED, or HEAD HTTP request.

4. The HTML code is cleaned and normalized.

5. The character-set is detected and normalized to UTF-8.

6. (a) The page (in any of the formats PDF, PostScript, MSWord, MSExcel, MSPow-
   erPoint, RTF and TeX/LaTeX) is converted to HTML or plain text by an
   external program.
   (b) Internal parsers handles HTML, plain text and images. This step extracts
   structured information like metadata (title, keywords, description ...), HTML
   links, and text without markup.

---

[15]http://www.robotstxt.org/wc/exclusion.html

7. The document is sent to the topic filter (see section **??**). If the Web-page is relevant with respect to the focus topic, processing continues with:

   (a) Heuristics like score propagation.
   (b) Further analysis, like genre and language identification.
   (c) Updating the record database.
   (d) Updating the frontier database with HTML links and URLs extracted from plain text.

Depending on several factors like configuration, hardware, network, workload, the crawler normally processes between 50 and 200 URLs per minute.

4.1   URL selection criteria

In order to successfully select and crawl one URL the following conditions (in this order) have to be met:

1. The URL has to be selected by the scheduling algorithm (section **??**).

   Relevant configuration variables: WaitIntervalHost (section **??**), WaitIntervalHarvesterLockRobotRules (section **??**), WaitIntervalHarvesterLockSuccess (section **??**)

2. The URL has to pass the allow test.
   Relevant configuration variables: allow (section **??**)

3. The URL is not be excluded by the exclude test (see section **??**).
   Relevant configuration variables: exclude (section **??**)

4. The Robot Exclusion Protocol has to allow crawling of the URL.

5. Optionally the document at the URL location has to pass the topic filter (section **??**).
   Relevant configuration variables: classifyPlugIn (section **??**), doCheckRecord (section **??**).

4.2   Document parsing and information extraction

Each document is parsed and analyzed by the crawler in order to store structured document records in the internal MySQL database. The structure of the record includes the fields:

- Title
- Headings
- Metadata
- Plain text
- Original document
- Links – HTML and plain text URLs

14

- Link anchor text

- Mime-Type

- Dates – modification, expire, and last checked by crawler

- Web-server identification

Optional some extra analysis can be done, see section **??**.

The system selects a document parser based on the Mime-Type together with available parsers and converter programs.

1. For some mime-types an external program is called in order to convert the document to a format handled internally (HTML or plain text).

   Relevant configuration variables: converters (section **??**)

2. Internal parsers handle HTML, plain text, TeX, and Image.

   Relevant configuration variables: converters (section **??**)

Supporting a new document format is as easy as providing a program that can convert a document in this format to HTML or plain text. Configuration of the mapping between document format (Mime-Type) and converter program is done in the complex configuration variable 'converters' (section **??**).

Out of the box Combine handle the following document formats: plain text, HTML, PDF, PostScript, MSWord, MSPowerPoint, MSExcel, RTF, TeX, and images.

4.3 URL filtering

Before an URL is accepted for scheduling (either by manual loading or recycling) it is normalized and validated. This process comprises a number of steps:

- Normalization

  - General practice: host-name lowercasing, port-number substitution, canonical URL
  - Removing fragments (ie '#' and everything after that)
  - Cleaning CGI repetitions of parameters
  - Collapsing dots ('./', '../') in the path
  - Removing CGI parameters that are session ids, as identified by patterns in the configuration variable sessionids (section **??**)
  - Normalizing Web-server names by resolving aliases. Identified by patterns in the configuration variable serveralias (section **??**). These patterns can be generated by using the program `combineUtil` to analyze a crawled corpus.

- Validation: A URL has to pass all three validation steps outlined below.

  - URL length has to be less than configuration variable maxUrlLength (section **??**)
  - Allow test: one of the Perl regular expressions in the configuration variable allow (section **??**) must match the URL

– Exclude test: none of the Perl regular expressions in the configuration variable exclude (section **??**) must match the URL

Both allow and exclude can contain two types of regular expressions identified by either 'HOST:' or 'URL' in front of the regular expression. The 'HOST:' regular expressions are matched only against the Web-server part of the URL while the 'URL' regular expressions are matched against the entire URL.

4.4 Crawling strategy

The crawler is designed to run continuously in order to keep crawled databases as up-to-date as possible. Starting and halting crawling is done manually. The configuration variable AutoRecycleLinks (section **??**) determines if the crawler should follow all valid new links or just take those that already are marked for crawling.

All links from a relevant document are extracted, normalized and stored in the structured record. Those links that pass the selection/validation criteria outlined below are marked for crawling.
To mark a URL for crawling requires:

- The URL should be from a page that is relevant (i.e. pass the focus filter).

- The URL scheme must be one of HTTP, HTTPS, or FTP.

- The URL must not exceed the maximum length (configurable, default 250 characters).

- It should pass the 'allow' test (configurable, default all URLs passes).

- It should pass the 'exclude' test (configurable, default excludes malformed URLs, some CGI pages, and URLs with file-extensions for binary formats).

At each scheduling point one URL from each available (unlocked) host is selected to generate a ready queue, which is then processed completely before a new scheduling is done. Each selected URL in the ready queue thus fulfills these requirements:

- URL must be marked for crawling (see above).

- URL must be unlocked (each successful access to a URL locks it for a configurable time WaitIntervalHarvesterLockSuccess (section **??**)).

- Host of the URL must be unlocked (each access to a host locks it for a configurable time WaitIntervalHost (section **??**)).

This implements a variant of BreathFirst crawling where a page is fetched if and only if a certain time threshold is exceeded since the last access to the server of that page.

4.5 Built-in topic filter – automated subject classification using string matching

The built-in topic filter is an approach to automated classification, that uses a topic definition with a pre-defined controlled vocabulary of topical terms, to determine relevance judgement. Thus it does not rely on a particular set of seed pages, or a collection of pre-classified example pages to learn from. It does require that some of the seed pages are

relevant and contain links into the topical area. One simple way of creating a set of seed pages would be to use terms from the controlled vocabulary as queries for a general-purpose search engine and take the result as seed pages.

The system for automated topic classification (overview in Figure **??**), that determines topical relevance in the topical filter, is based on matching subject terms from a controlled vocabulary in a topic definition with the text of the document to be classified [**?**]. The topic definition uses subject classes in a hierarchical classification system (corresponding to topics) and terms associated with each subject class. Terms can be single words, phrases, or Boolean AND-expressions connecting terms. Boolean OR-expressions are implicitly handled by having several different terms associated with the same subject class (see section **??**).

The algorithm works by string-to-string matching of terms and text in documents. Each time a match is found the document is awarded points based on which term is matched and in which structural part of the document (location) the match is found [**?**]. The points are summed to make the final relevance score of the document. If the score is above a cut-off value the document is saved in the database together with a (list of) subject classification(s) and term(s).

Figure 3: Overview of the automated topic classification algorithm

By providing a list of known relevant sites in the configuration file `sitesOK.txt` (located in the job specific configuration directory) the above test can be bypassed. It works by checking the host part of the URL against the list of known relevant sites and if a match is found the page is validated and saved in the database regardless of the outcome of the algorithm.

4.5.1 Topic definition

Located in `/etc/combine/<jobname>/topicdefinition.txt`. Topic definitions use triplets (term, relevance weight, topic-classes) as its basic entities. Weights are signed integers and indicate the relevance of the term with respect to the topic-classes. Higher values indicate

more relevant terms. A large negative value can be used to exclude documents containing that term. Terms should be all lowercase.

Terms can be:

- single words

- a phrase (i.e. all words in exact order)

- a Boolean AND-expression connecting terms (i.e. all terms must be present but in any order). The Boolean AND operator is encoded as '@and'.

A Boolean OR-expression has to be entered as separate term triplets. The Boolean expression "`polymer AND (atactic OR syndiotactic)`" thus has to be translated into two separate triplets, one containing the term "`polymer @and atactic`", and another with "`polymer @and syndiotactic`".

Terms can include (Perl) regular expressions like:

- a '?' makes the character immediately preceding optional, i.e. the term "`coins?`" will match both "`coin`" and "`coins`"

- a "`[^\s]*`" is truncation (matches all character sequences except space ' '), "`glass art[^\s]*`" will match "`glass art`", "`glass arts`", "`glass artists`", "`glass articles`", and so on.

It is important to understand that each triplet in the topic definition is considered by itself without any context, so they must **each** be topic- or sub-class specific in order to be useful. Subject neutral terms like "use", "test", "history" should not be used. If really needed they have to be qualified so that they become topic specific (see examples below). Simple guidelines for creating the triplets and assigning weights are:

- Phrases or unique, topic-specific terms, should be used if possible, and assigned the highest weights, since they normally are most discriminatory.

- Boolean AND-expressions are the next best.

- Single words can be too general and/or have several meanings or uses that make them less specific and those should thus be assigned a small weights.

- Acronyms can be used as terms if they are unique.

- Negative weights should be used in order to exclude concepts.

4.5.2 Topic definition (term triplets) BNF grammar

TERM-LIST :== TERM-ROW '<**cr**>' || '#' <**char**>+ '<**cr**>' || '<**cr** >'
TERM-ROW :== WEIGHT ': ' TERMS '=' CLASS-LIST
WEIGHT :== ['-']<**integer**>
TERMS :== TERM [' @and ' TERMS]*
TERM :== WORD ' ' [WORD]*
WORD :== <**char**>+||<**char**>+<**perl-reg-exp**>
CLASS-LIST :== CLASSID [',' CLASS-LIST]

CLASSID :== **<char>+**

A line that starts with '#' is ignored and so are empty lines.

<perl-reg-exp> is only supported by the plain matching algorithm described in section **??**.

"CLASSID" is a topic (sub-)class specifier, often from a hierarchical classification system like Engineering Index[16].

4.5.3 Term triplet examples

```
50: optical glass=A.14.5, D.2.2
30: glass @and fiberoptics=D.2.2.8
50: glass @and technical @and history=D.2
50: ceramic materials @and glass=D.2.1.7
-10000: glass @and art=A
```

The first line says that a document containing the term "`optical glass`" should be awarded 50 points for each of the two classes A.14.5 and D.2.2.

"glass" as a single term is probably too general, qualify it with more terms like: "`glass @and fiberoptics`", or "`glass @and technical @and history`" or use a phrase like "`glass fiber`" or "`optical glass`".

In order to exclude documents about artistic use of glass the term "`glass @and art`" can be used with a (high) negative score.

An example from the topic definition for 'Carnivorous Plants' using regular expressions is given below:

```
#This is a comment
75: d\.?\s*californica=CP.Drosophyllum
10: pitcher[^\s]*=CP
-10: pitcher[^\s]* @and baseball=CP
```

The term "`d\.?\s*californica`" will match `D californica`, `D. californica`, `D.californica` etc.

The last two lines assure that a document containing "`pitcher`" gets 10 points but if the document also contains "`baseball`" the points are removed.

4.5.4 Algorithm 1: plain matching

This algorithm is selected by setting the configuration parameter

   `classifyPlugIn = Combine::Check_record`

The algorithm produces a list of suggested topic-classes (subject classifications) and corresponding relevance scores using the algorithm:

$$\text{Relevance\_score} =$$

$$\sum_{\text{all locations}} \left(\sum_{\text{all terms}} (hits[\text{location}_j][\text{term }_i] * weight[\text{term}_i] * weight[\text{location}_j]) \right)$$

---

[16]http://www.ei.org/

term weight ($weight[\text{term}_i]$) is taken from the topic definition triplets.

location weight ($weight[\text{location}_j]$) are defined ad hoc for locations like title, metadata, HTML headings, and plain text. However the exact values for these weights does not seem to play a large role in the precision of the algorithm [**?**].

hits ($hits[\text{location}_j][\text{term}_i]$) is the number of times term_i occur in the text of location_j

The summed relevance score might, for certain applications, have to be normalized with respect to text size of the document.

One problem with this algorithm is that a term that is found in the beginning of the text contributes as much as a term that is found at the end of a large document. Another problem is the distance and thus the coupling between two terms in a Boolean expression might be very large in a big document and this is not taken into account by the above algorithm.

4.5.5 Algorithm 2: position weighted matching

This algorithm is selected by setting the configuration parameter

```
classifyPlugIn = Combine::PosCheck_record
```

In response to the problems cited above we developed a modified version of the algorithm that takes into account word position in the text and proximity for Boolean terms. It also eliminates the need to assign ad hoc weights to locations. The new algorithm works as follows.

First all text from all locations are concatenated (in the natural importance order title, metadata, text) into one chunk of text. Matching of terms is done against this chunk. Relevance score is calculated as

$$\text{Relevance\_score} =$$

$$\sum_{\text{all terms}} \left(\sum_{\text{all matches}} \frac{weight[\text{term}_i]}{\log(k * position[\text{term}_i][\text{match}_j]) * proximity[\text{term}_i][\text{match}_j]} \right)$$

term weight ($weight[\text{term}_i]$) is taken from the topic definition triplets

position ($position[\text{term}_i][\text{match}_j]$) is the position in the text (starting from 1) for match_j of term_i. The constant factor k is normally 0.5

proximity ($proximity[\text{term}_i][\text{match}_j]$) is

$$1 \qquad \text{for non Boolean terms}$$
$$\log(distance\_between\_components) \quad \text{for Boolean terms}$$

In this algorithm a matched term close to the start of text contributes more to the relevance score than a match towards the end of the text. And for Boolean terms the closer the components are the higher the contribution to the relevance score.

4.6 Built-in topic filter – automated subject classification using SVM

Topic filetring using SVM (Support Vector Machines) classifiers are supported using the SVMLight package[17]. This package has to be installed manually together with the Algo-

---

[17] http://svmlight.joachims.org/

rithm::SVMLight Perl module. For installation hints see CPAN SVMLight README[18] or 'installing-algorithm-svmlight-linux-ubuntu[19],'

SVM classifiers need a trained model before they can be used.

The procedure to get started is as follows:

- Make sure that Algorithm::SVMLight[20] and SVMLight[21] are installed.

- Collect examples of good and bad URLs that defines your topic (the more the better).

- Generate a SVM model with the program `combineSVM`.

- Initialize a new job with `combineINIT`.

- Copy the SVM model to the job's configuration directory `/etc/combine/<jobname>/SVMmodel.txt`.

- Edit the configuration file `/etc/combine/<jobname>/combine.cfg` and add the following:

```
doCheckRecord = 1
classifyPlugIn = Combine::classifySVM
SVMmodel = SVMmodel.txt
```

- Then proceed with crawling as normal.

4.7   Topic filter Plug-In API

The configuration variable classifyPlugIn (section **??**) is used to find the Perl module that implements the desired topic filter. The value should be formatted as a valid Perl module identifier (i.e. the module must be somewhere in the Perl module search path). Combine will call a subroutine named 'classify' in this module, providing an XWI-object as in parameter. An XWI-object is a structured object holding all information from parsing a Web-page. The subroutine must return either 0 or 1, where

0: means record fails to meet the classification criteria, i.e. ignore this record

1: means record is OK, store it in the database, and follow the links

More details on how to write a Plug-In can be found in the example classifyPlugIn-Template.pm (see Appendix **??**).

4.8   Analysis

Extra analysis, enabled by the configuration variable doAnalyse (section **??**), tries to determine the language of the content and the country of the Web-server. Both are stored in the internal database.

---

[18]http://search.cpan.org/src/KWILLIAMS/Algorithm-SVMLight-0.08/README
[19]http://www.dale-emmons.com/content/installing-algorithm-svmlight-linux-ubuntu
[20]http://search.cpan.org/ kwilliams/Algorithm-SVMLight-0.08/
[21]http://svmlight.joachims.org/

4.9 Duplicate detection

Duplicates of crawled documents are automatically detected with the aid of a MD5-checksum calculated on the contents of the document.

The MD5-checksum is used as the master record key in the internal database thus preventing pollution with duplicate pages. All URLs for a page are stored in the record, and a page is not deleted from the database until the crawler has verified that it's unavailable from all the saved URLs.

4.10 URL recycling

URLs for recycling come from 3 sources:

- Links extracted during HTML parsing.

- Redirects (unless configuration variable UserAgentFollowRedirects (section **??**) is set).

- URLs extracted from plain text (enabled by the configuration variable extractLinks-FromText (section **??**)).

Automatic recycling of URLs is enabled by the configuration variable AutoRecycleLinks (section **??**). It can also be done manually with the command
`combineCtrl --jobname XXXX recyclelinks`
The command `combineCtrl --jobname XXXX reharvest` marks all pages in the database for harvesting again.

4.11 Database cleaning

The tool `combineUtil` implements functionality for cleaning the database.

sanity/restoreSanity checks respectively restore consistency of the internal database.

deleteNetLoc/deletePath/deleteMD5/deleteRecordid deletes records from the database based on supplied parameters.

serverAlias detects Web-server aliases in the database. All detected alias groups are added to the serveralias configuration (section **??**). Records from aliased servers (except for the first Web-server) will be deleted.

4.12 Complete application – SearchEngine in a Box

The SearchEngine-in-a-Box[22] system is based on the two systems Combine Focused Crawler and Zebra text indexing and retrieval engine[23]. This system allows you build a vertical search engine for your favorite topic in a few easy steps.

The SearchEngine-in-a-Box Web-site contains instructions and downloads to make this happen. Basically it makes use of the ZebraHost (see section **??**) configuration variable which enables direct communication between the crawler and the database system and thus indexes records as soon as they are crawled. This also means that they are directly searchable.

---

[22]http://combine.it.lth.se/SearchEngineBox/

[23]http://www.indexdata.dk/zebra/

5 Evaluation of automated subject classification

5.1 Approaches to automated classification

According to [?], one can distinguish between three major approaches to automated classification: text categorization, document clustering, and string-to-string matching.

Machine learning, or text categorization, is the most widespread approach to automated classification of text, in which characteristics of pre-defined categories are learnt from intellectually categorized documents. However, intellectually categorized documents are not available in many subject areas, for different types of documents or for different user groups. For example, today the standard text classification benchmark is a Reuters RCV1 collection [?], which has about 100 classes and 800000 documents. This would imply that for a text categorization task some 8000 training and testing documents per class are needed. Another problem is that the algorithm works only for that document collection on which parts it has been trained. In addition, [?] claims that the most serious problem in text categorization evaluations is the lack of standard data collections and shows how different versions of the same collection have a strong impact on performance, whereas other versions do not.

In document clustering, the predefined categories (the controlled vocabulary) are automatically produced: both clusters' labels and relationships between them are automatically generated. Labelling of the clusters is a major research problem, with relationships between the categories, such as those of equivalence, related-term and hierarchical relationships, being even more difficult to automatically derive ([?], p.168). "Automatically-derived structures often result in heterogeneous criteria for category membership and can be difficult to understand" [?]. Also, clusters change as new documents are added to the collection.

In string-to-string matching, matching is conducted between a controlled vocabulary and text of documents to be classified. This approach does not require training documents. Usually weighting schemes are applied to indicate the degree to which a term from a document to be classified is significant for the document's topicality. The importance of the controlled vocabularies such as thesauri in automated classification has been recognized in recent research. [?] used a thesaurus to improve performance of the k-NN classifier and managed to improve precision for about 14 %, without degrading recall. [?] showed how information from a subject-specific thesaurus improved performance of key-phrase extraction by more than 1,5 times in F1, precision, and recall. [?] demonstrated that subject ontologies could help improve word sense disambiguation.

Thus, the chosen approach to automated subject classification in the crawler was string-matching. Apart from the fact that no training documents are required, a major motivation to apply this approach was to re-use the intellectual effort that has gone into creating such a controlled vocabulary. Vocabulary control in thesauri is achieved in several ways, out of which the following are beneficial for automated classification:

- the terms are usually noun phrases, which are content words;

- the meaning of the term is restricted to that most effective for the purposes of a particular thesaurus, which is indicated by the addition of scope notes and definitions, providing additional context for automated classification;

- three main types of relationships are displayed in a thesaurus: 1) equivalence (synonyms, lexical variants, terms treated as synonyms for general purposes); 2) hierarchical (generic, whole-part or instance relationships); 3) associative (terms that

are closely related conceptually but not hierarchically and are not members of an equivalence set).

In automated classification, equivalence terms allow for discovering the concepts and not just words expressing them. Hierarchies provide additional context for determining the correct sense of a term, and so do associative relationships.

5.1.1 Description of the used string-matching algorithm

Automated classification approach used for evaluation was string-matching of terms (cf. section **??**) from an engineering-specific controlled vocabulary Engineering Index (Ei) thesaurus and classification scheme, used in Elsevier's Compendex database. Ei classification scheme is organized into six categories which are divided into 38 subjects, which are further subdivided into 182 specific subject areas. These are further subdivided, resulting in some 800 individual classes in a five-level hierarchy. In Ei there are on average 88 intellectually selected terms designating one class.

The algorithm searches for terms from the Ei thesaurus and classification scheme in documents to be classified. In order to do this, a term list is created, containing class captions, different thesauri terms and classes which the terms and captions denote. The term list consists of triplets: term (single word, Boolean term or phrase), class which the term designates or maps to, and weight. Boolean terms consist of words that must all be present but in any order or in any distance from each other. The Boolean terms are not explicitly part of the Ei thesaurus, so they had to be created in a pre-processing step. They are considered to be those terms which contain the following strings: 'and', 'vs.' (short for versus), ',' (comma), ';' (semi-colon, separating different concepts in class names), '(' (parenthesis, indicating the context of a homonym), ':' (colon, indicating a more specific description of the previous term in a class name), and '--' (double dash, indicating heading–subheading relationship). Upper-case words from the Ei thesaurus and classification scheme are left in upper case in the term list, assuming that they are acronyms. All other words containing at least one lower-case letter are converted into lower case. Geographical names are excluded on the grounds that they are not being engineering-specific in any sense.

The following is an excerpt from the Ei thesaurus and classification scheme, based on which the excerpt from the term list (further below) was created.
From the classification scheme (captions):

```
931.2 Physical Properties of Gases, Liquids and Solids
...
942.1 Electric and Electronic Instruments
...
943.2 Mechanical Variables Measurements
```

From the thesaurus:

```
TM Amperometric sensors
UF Sensors--Amperometric measurements
MC 942.1

TM Angle measurement
UF Angular measurement
```

```
UF Mechanical variables measurement--Angles
BT Spatial variables measurement
RT Micrometers
MC 943.2

TM Anisotropy
NT Magnetic anisotropy
MC 931.2
```

TM stands for the preferred term, UF for synonym, BT for broader term, RT for related term, NT for narrower term; MC represents the main class. Below is an excerpt from one term list, as based on the above examples:

```
1: electric @and electronic instruments=942.1,
1: mechanical variables measurements=943.2,
1: physical properties of gases @and liquids @and solids=931.2,
1: amperometric sensors=942.1,
1: sensors @and amperometric measurements=942.1,
1: angle measurement=943.2,
1: angular measurement=943.2,
1: mechanical variables measurement @and angles=943.2,
1: spatial variables measurement=943.2,
1: micrometers=943.2,
1: anisotropy=931.2,
1: magnetic anisotropy=931.2
```

The algorithm looks for strings from a given term list in the document to be classified and if the string (e.g. 'magnetic anisotropy' from the above list) is found, the class(es) assigned to that string in the term list ('931.2' in our example) are assigned to the document. One class can be designated by many terms, and each time the class is found, the corresponding weight ('1' in our example) is assigned to the class.

The scores for each class are summed up and classes with scores above a certain cut-off (heuristically defined) can be selected as the final ones for that document. Experiments with different weights and cut-offs are described in the following sections.

5.2 Evaluation methodology

5.2.1 Evaluation challenge

According to [?], intellectually-based subject indexing is a process involving the following three steps: determining the subject content of the document, conceptual analysis to decide which aspects of the subject should be represented, and translation of those concepts or aspects into a controlled vocabulary. These steps are based on the library's policy in respect to their collections and user groups. Thus, when in automated classification study, the used document collection is the one in which documents have been indexed intellectually, the policies based on which indexing has been conducted should also be known, and automated classification should then be based on those policies as well.

Ei thesaurus and classification scheme is rather big and deep (five hierarchical levels), allowing many different choices. Without a thorough qualitative analysis of automatically assigned classes one cannot be sure if, for example, the classes assigned by the algorithm,

which were not intellectually assigned, are actually wrong, or if they were left-out by mistake or because of the indexing policy.

In addition, subject indexers make errors such as those related to exhaustivity policy (too many or too few terms get assigned), specificity of indexing (usually this error means that not the most specific term found was assigned), they may omit important terms, or assign an obviously incorrect term ([?], p.86-87). In addition, it has been reported that different people, whether users or subject indexers, would assign different subject terms or classes to the same document. Studies on inter-indexer and intra-indexer consistency report generally low indexer consistency ([?], p. 99-101). There are two main factors that seem to affect it: 1) higher exhaustivity and specificity of subject indexing both lead to lower consistency (indexers choose the same first term for the major subject of the document, but the consistency decreases as they choose more classes or terms); 2) the bigger the vocabulary, or, the more choices the indexers have, the less likely they will choose the same classes or terms (ibid.). Few studies have been conducted as to why indexers disagree [?].

Automated classification experiments today are mostly conducted under controlled conditions, ignoring the fact that the purpose of automated classification is improved information retrieval, which should be evaluated in context (cf. [?]). As Sebastiani ([?] p. 32) puts it, "the evaluation of document classifiers is typically conducted experimentally, rather than analytically. The reason is that we would need a formal specification of the problem that the system is trying to solve (e.g. with respect to what correctness and completeness are defined), and the central notion that of membership of a document in a category is, due to its subjective character, inherently nonformalizable".

Due to the fact that methodology for such experiments has yet to be developed, as well as due to limited resources, we follow the traditional approach to evaluation and start from the assumption that intellectually assigned classes in the data collection are correct, and the results of automated classification are being compared against them.

5.2.2 Evaluation measures used

Assuming that intellectually assigned classes in the data collection are correct, evaluation in this study is based on comparison of automatically derived classes against the intellectually assigned ones. Ei classes are much related to each other and often there is only a small topical difference between them. The topical relatedness is expressed in numbers representing the classes: the more initial digits any two classes have in common, the more related they are. Thus, comparing the classes at only the first few digits instead of all the five (each representing one hierarchical level), would also make sense. Evaluation measures used were the standard microaveraged and macroaveraged precision, recall and F1 measures ([?], p.33) for complete matching of all digits as well as for partial matching.

5.2.3 Data collection

In the experiments, the following data collections were used:

1. For deriving significance indicators assigned to different Web page elements [?], and identifying issues specific to Web pages [?] some 1000 Web pages engineering, to which Ei classes had been manually assigned as part of the EELS subject gateway[24] were used.

---

[24]http://eels.lub.lu.se/

26

2. For deriving weights based on the type of controlled vocabulary term [**?**], and for enriching the term list with terms extracted using natural language processing, the data collection consisted of 35166 document records from the Compendex database[25]. From each document record the following elements were utilized: an identification number, title, abstract and intellectually pre-assigned classes, for example:

Identification number: 03337590709
Title: `The concept of relevance in IR`
Abstract: `This article introduces the concept of relevance as viewed and applied in the context of IR evaluation, by presenting an overview of the multidimensional and dynamic nature of the concept. The literature on relevance reveals how the relevance concept, especially in regard to the multidimensionality of relevance, is many faceted, and does not just refer to the various relevance criteria users may apply in the process of judging relevance of retrieved information objects. From our point of view, the multidimensionality of relevance explains why some will argue that no consensus has been reached on the relevance concept. Thus, the objective of this article is to present an overview of the many different views and ways by which the concept of relevance is used - leading to a consistent and compatible understanding of the concept. In addition, special attention is paid to the type of situational relevance. Many researchers perceive situational relevance as the most realistic type of user relevance, and therefore situational relevance is discussed with reference to its potential dynamic nature, and as a requirement for interactive information retrieval (IIR) evaluation.`
Ei classification codes: `903.3 Information Retrieval & Use, 723.5 Computer Applications, 921 Applied Mathematics`

In our collection we included only those documents that have at least one class in the area of Engineering, General, covered by 92 classes we selected. The subset of 35166 documents was selected from the Compendex database by simply retrieving the first documents offered by the Compendex user interface, without changing any preferences. The query was to find those documents that were assigned a certain class. A minimum of 100 documents per class was retrieved at several different points in time during the last year. Compendex is a commercial database so the subset cannot be made available to others. However, the authors can provide documents' identification numbers on request. In the data collection there were on average 838 documents per class, ranging from 138 to 5230.

3. For comparing classification performance of the string-matching algorithm against a machine-learning one [**?**], the data collection consisted of a subset of paper records from the Compendex database, classified into six selected classes. In this run of the experiment, only the six classes were selected in order to provide us with indications for further possibilities. Classes 723.1.1 (Computer Programming Languages), 723.4 (Artificial Intelligence), and 903.3 (Information Retrieval and Use) each had 4400 examples (the maximum allowed by the Compendex database provider), 722.3 (Data Communication Equipment and Techniques) 2800, 402 (Buildings and Towers) 4283, and 903 (Information Science) 3823 examples.

---

[25]http://www.engineeringvillage2.org/

5.3 Results

5.3.1 The role of different thesauri terms

In one study [?], it has been explored to what degree different types of terms in Engineering Index influence automated subject classification performance. Preferred terms, their synonyms, broader, narrower, related terms, and captions were examined in combination with a stemmer and a stop-word list. The best performance measured as mean F1 macroaveraged and microaverage F1 values has the preferred term list, and the worst one the captions list. Stemming showed to be beneficial in four out of the seven different term lists: captions, narrower, preferred, and synonyms. Concerning stop words, the mean F1 improved for narrower and preferred terms. Concerning the number of classes per document that get automatically assigned, when using captions less than one class is assigned on average even when stemming is applied; narrower and synonyms improve with stemming, close to our aim of 2,2 classes that have been intellectually assigned. The most appropriate number of classes get assigned when preferred terms are used with stop-words.

Based on other term types, too many classes get assigned, but that could be dealt with in the future by introducing cut-offs. Each class is on average designated by 88 terms, ranging from 1 to 756 terms per class. The majority of terms are related terms, followed by synonyms and preferred terms. By looking at the 10 top-performing classes, it was shown that the sole number of terms designating a class does not seem to be proportional to the performance. Moreover, these best performing classes do not have a similar distribution of types of terms designating them, i.e. the percentage of certain term types does not seem to be directly related to performance. The same was discovered for the 10 worst-performing classes.

In conclusion, the results showed that preferred terms perform best, whereas captions perform worst. Stemming in most cases showed to improve performance, whereas the stop-word list did not have a significant impact. The majority of classes is found when using all the terms and stemming: micro-averaged recall is 73 %. The remaining 27 % of classes were not found because the words in the term list designating the classes did not exist in the text of the documents to be classified. This study implies that all types of terms should be used for a term list in order to achieve best recall, but that higher weights could be given to preferred terms, captions and synonyms, as the latter yield highest precision. Stemming seems useful for achieving higher recall, and could be balanced by introducing weights for stemmed terms. Stop-word list could be applied to captions, narrower and preferred terms.

5.3.2 Enriching the term list using natural language processing

In order to allow for better recall, the basic term list was enriched with new terms. From the basic term list, preferred and synonymous terms were taken, since they give best precision, and based on them new terms were extracted. These new terms were derived from documents issued from the Compendex database, using multi-word morpho-syntactic analysis and synonym acquisition. Multi-word morpho-syntactic analysis was conducted using a unification-based partial parser FASTER which analyses raw technical texts and, based on built-in meta-rules, detects morpho-syntactic variants. The parser exploits morphological (derivational and inflectional) information as given by the database CELEX. Morphological analysis was used to identify derivational variants (e.g. effect of gravity: gravitational effect), and syntactical analysis to insert word inside a term (e.g. flow measurement: flow

discharge measurements), permute components of a term (e.g. control of the inventory: inventory control) or add a coordinated component to a term (e.g. control system: control and navigation system).

Synonyms were acquired using a rule-based system, SynoTerm which infers synonymy relations between complex terms by employing semantic information extracted from lexical resources. Documents were first preprocessed and tagged with part-of-speech information and lemmatized. Terms were then identified using the term extractor YaTeA based on parsing patterns and endogenous disambiguation. The semantic information provided by the database WordNet was used as a bootstrap to acquire synonym terms of the basic terms.

The number of classes that were enriched using these natural language processing methods is as follows: derivation 705, out of which 93 adjective to noun, 78 noun to adjective, and 534 noun to verb derivations; permutation 1373; coordination 483; insertion 742; preposition change 69; synonymy 292 automatically extracted, out of which 168 were manually verified as correct.

By combining all the extracted terms into one term list, the mean F1 is 0.14 when stemming is applied, and microaveraged recall is 0.11. This implies that enriching the original Ei-based term list should improve recall. In comparison to results we get when gained with the original term list (micro-averaged recall with stemming 0.73), here the best recall, also microaveraged and with stemming, is 0.76.

5.3.3 Importance of HTML structural elements and metadata

In [?] the aim was to determine how significance indicators assigned to different Web page elements (internal metadata, title, headings, and main text) influence automated classification. The data collection that was used comprised 1000 Web pages in engineering, to which Engineering Information classes had been manually assigned. The significance indicators were derived using several different methods: (total and partial) precision and recall, semantic distance and multiple regression. It was shown that for best results all the elements have to be included in the classification process. The exact way of combining the significance indicators turned out not to be overly important: using the F1 measure, the best combination of significance indicators yielded no more than 3 % higher performance results than the baseline.

5.3.4 Challenges and recommendations for classification of Web pages

Issues specific to Web pages were identified and discussed in [?]. The focus of the study was a collection of Web pages in the field of engineering. Web pages present a special challenge: because of their heterogeneity, one principle (e.g. words from headings are more important than main text) is not applicable to all the Web pages of a collection. For example, utilizing information from headings on all Web pages might not give improved results, since headings are sometimes used simply instead of using bold or a bigger font size.

A number of weaknesses of the described approach were identified:

1. Class not found at all, because the words in the term list designating the classes were not found in the text of the Web page to be classified.

2. Class found but below threshold, which has to do with weighting and cut-off values. This is because in the approach only classes with scores above a pre-defined cut-off

value are selected as the classes for the document: the final classes selected are those with scores that contain at least 5 % of the sum of all the scores assigned in total, or, if such a class doesn't exist, the class with the top score is selected. Another reason could be that the classification algorithm is made to always pick the most specific class as the final one, which is in accordance with the given policy for intellectual classification.

3. Wrong automatically assigned class. Based on the sample, four different sub-problems have been identified:

   - words recognized as homonyms or distant synonyms;
   - word found on a Web page is there because it is an instance of what it represents, and it is not about such instances (e.g. a Web page on online tutorials and e-learning programs for technical fields gets wrongly classified as a Web page on 'education');
   - too distant term-class mappings, including cases when one term in the term list is mapped to several different classes;
   - words mentioned on the Web page have little to do with the Web page's aboutness, e.g. an institution's Web page gets wrongly classified as 'facsimile systems and technology', because among their contact information, there is also a fax number, and the word 'fax' is mapped to that class.

4. Automatically assigned class that is not really wrong, which probably has to do with the subject indexing policy such as exhaustivity.

Ways to deal with those problems were proposed for further research. These include enriching the term list with synonyms and different word forms, adjusting the term weights and cut-off values and word-sense disambiguation. In our further research the plan is to implement automated methods. On the other hand, the suggested manual methods (e.g. adding synonyms) would, at the same time, improve Ei's original function, that of enhancing retrieval. Having this purpose in mind, manually enriching a controlled vocabulary for automated classification or indexing would not necessarily create additional costs.

5.3.5 Comparing and combining two approaches

In [?] a machine-learning and a string-matching approach to automated subject classification of text were compared as to their performance on a test collection of six classes. Our first hypothesis was that, as the string-matching algorithm uses manually constructed model, we expect it to have higher precision than the machine learning with its automatically constructed model. On the other hand, while the machine-learning algorithm builds the model from the training data, we expect it to have higher recall in addition to being more flexible to changes in the data. Experiments have confirmed the hypothesis only on one of the six classes. Experimental results showed that SVM on average outperforms the string-matching algorithm. Different results were gained for different classes. The best results in string-matching are for class 402, which we attribute to the fact that it has the highest number of term entries designating it (423). Class 903.3, on the other hand, has only 26 different term entries designating it in the string-matching term list, but string-matching largely outperforms SVM in precision (0.97 vs. 0.79). This is subject to further investigation.

The two approaches being complementary, we investigated different combinations of the two based on combining their vocabularies. The linear SVM in the original setting was trained with no feature selection except the stop-word removal. Additionally, three experiments were conducted using feature selection, taking:

1. only the terms that are present in the controlled vocabulary;

2. the top 1000 terms from centroid tf-idf vectors for each class (terms that are characteristic for the class - descriptive terms);

3. the top 1000 terms from the SVM-normal trained on a binary classification problem for each class (terms that distinguish one class form the rest distinctive terms).

In the experiments with string-matching algorithm, four different term lists were created, and we report performance for each of them:

1. the original one, based on the controlled vocabulary;

2. the one based on automatically extracted descriptive keywords from the documents belonging to their classes;

3. the one based on automatically extracted distinctive keywords from the documents belonging to their classes;

4. the one based on union of the first and the second list.

SVM performs best using the original set of terms, and string-matching approach also has best precision when using the original set of terms. Best recall for string-matching is achieved when using descriptive terms. Reasons for these results need further investigation, including a larger data collection and combining the two using predictions.

6   Performance and scalability

Performance evaluation of the automated subject classification component is treated in section **??**.

6.1   Speed

Performance in terms of number of URLs treated per minute is of course highly dependent on a number of circumstances like network load, capacity of the machine, the selection of URLs to crawl, configuration details, number of crawlers used, etc. In general, within rather wide limits, you could expect the Combine system to handle up to 200 URLs per minute. By "handle" we mean everything from scheduling of URLs, fetching pages over the network, parsing the page, automated subject classification, recycling of new links, to storing the structured record in a relational database. This holds for small simple crawls starting from scratch to large complicated topic specific crawls with millions of records.

The prime way of increasing performance is to use more than one crawler for a job. This is handled by the `--harvesters` switch used together with the `combineCtrl start` command for example `combineCtrl --jobname MyCrawl --harvesters 5 start` will start 5 crawlers working together on the job 'MyCrawl'. The effect of using more than one crawler on crawling speed is illustrated in figure **??** and the resulting speedup is shown in table **??**.

Figure 4: Combine crawler performance, using no focus and configuration optimized for speed.

No of crawlers	1	2	5	10	15	20
Speedup	1	2.0	4.8	8.2	9.8	11.0

Table 1: Speedup of crawling vs number of crawlers

Configuration also has an effect on performance. In Figure **??** performance improvements based on configuration changes are shown. The choice of algorithm for automated classification turns out to have biggest influence on performance, where algorithm 2 – section **??** – (`classifyPlugIn = Combine::PosCheck_record` – Pos in Figure **??**) is much faster than algorithm 1 – section **??** – (`classifyPlugIn = Combine::Check_record` – Std in Figure **??**). Configuration optimization consisted of not using Tidy to clean HTML (`useTidy = 0`) and not storing the original page in the database (`saveHTML = 0`). Tweaking of other configuration variables (like disabling logging to the MySQL database `Loglev = 0`) also has an effect on performance but to a lesser degree.

6.2 Space

Storing structured records including the original document takes quite a lot of disk space. On average 25 kB per record is used by MySQL. This includes the administrative overhead needed for the operation of the crawler. A database with 100 000 records needs at least 2.5 GB on disk. Deciding not to store the original page in the database (`saveHTML = 0`) gives considerable space savings. On average 8 kB per is used without the original HTML.

Exporting records in the ALVIS XML format further increases size to 42 kB per record.

Figure 5: Effect of configuration changes on focused crawler performance, using 10 crawlers and a topic definition with 2512 terms.

Using the slight less redundant XML-format `combine` uses 27 kB per record. Thus 100 000 records will generate a file of size 3 to 4 GB. The really compact Dublin Core format (`dc`) generates 0.65 kB per record.

6.3   Crawling strategy

In [?] four different crawling strategies are studied:

BreadthFirst The simplest strategy for crawling. It does not utilize heuristics in deciding which URL to visit next. It uses the frontier as a FIFO queue, crawling links in the order in which they are encountered.

BestFirst The basic idea is that given a frontier of URLs, the best URL according to some estimation criterion is selected for crawling, using the frontier as a priority queue. In this implementation, the URL selection process is guided by the topic score of the source page as calculated by Combine.

PageRank The same as Best-First but ordered by PageRank calculated from the pages crawled so far.

BreadthFirstTime A version of BreadthFirst. It is based on the idea of not accessing the same server during a certain period of time in order not to overload servers. Thus, a page is fetched if and only if a certain time threshold is exceeded since the last access to the server of that page.

Results from a simulated crawl (figure **??** from [**?**]) show that at first PageRank performs best but BreadthFirstTime (which is used in Combine) prevails in the long run, although differences are small.

Figure 6: Total number of relevant pages visited

7   System components

All executables take a mandatory switch `--jobname` which is used to identify the particular crawl job you want as well as the job-specific configuration directory.

Briefly `combineINIT` is used to initialize SQL database and the job specific configuration directory. `combineCtrl` controls a Combine crawling job (start, stop, etc.) as well as printing some statistics. `combineExport` exports records in various XML formats and `combineUtil` provides various utility operations on the Combine database.

Detailed dependency information (section **??**) can be found in the 'Gory details' section.

In appendix (**??**) you'll find all the man-pages collected.

7.1   combineINIT

Creates a MySQL database, database tables and initializes it. If the database exists it is dropped and recreated. A job-specific configuration directory is created in `/etc/combine/` and populated with a default configuration file.

If a topic definition filename is given, focused crawling using this topic defintion is enabled per default. Otherwise focused crawling is disabled, and Combine works as a general crawler.

7.2 combineCtrl

Implements various control functionality to administer a crawling job, like starting and stopping crawlers, injecting URLs into the crawl queue, scheduling newly found links for crawling, controlling scheduling, etc. This is the preferred way of controling a crawl job.

7.3 combineUtil

Implements a number of utilities both for extracting information:

- Global statistics about the database

- matched terms from topic definition

- topic classes assigned to documents

and for database maintenance:

- sanity check and restoration

- deleting records specified by either Web-server, URL path, MD5 checksum, or internal record identifier

- server alias detection and managing

7.4 combineExport

Export of structured records is done according to one of three profiles: `alvis`, `dc`, or `combine`. `alvis` and `combine` are very similar XML formats where `combine` is more compact with less redundancy and `alvis` contains some more information. `dc` is XML-encoded Dublin Core data.

The `alvis` profile format is defined by the Alvis Enriched Document XML Schema[26].

For flexibility a switch `--xsltscript` adds the possibility to filter the output using a XSLT script. The script is fed a record according to the `combine` profile and the result is exported.

Switches `--pipehost` and `--pipeport` makes combineExport send it's output directly to an Alvis[27] pipeline reader instead of printing on stdout. This together with the switch `--incremental`, which just exports changes since the last invocation, provides an easy way of keeping an external system like Alvis or a Zebra[28] database updated.

7.5 Internal executables and Library modules

`combine` is the main crawling machine in the Combine system and `combineRun` starts, monitors and restarts `combine` crawling processes.

7.5.1 Library

Main, crawler-specific, library components are collected in the `Combine::` Perl namespace.

---

[26]http://www.miketaylor.org.uk/tmp/alvis/d3.1/enriched-document.xsd
[27]http://www.alvis.info
[28]http://www.indexdata.dk/zebra/

References

[1] Documentation - Methods for examining documents, determining their subjects, and selecting index terms. International Organization for Standardization, Standard 5963-1985.

[2] Lifeboat for knowledge organization: indexing theory. http://www.db.dk/bh/Lifeboat_KO/CONCEPTS/indexing_theory.htm.

[3] A. Ardö and T. Koch. Automatic classification applied to the full-text Internet documents in a robot-generated subject index. In *Online Information 99, Proceedings*, pages 239–246, Dec. 1999. http://www.it.lth.se/anders/online99/.

[4] S. L. Bang, J. D. Yang, and H. J. Yang. Hierarchical document categorization with k-nn and concept-based thesauri. *Information Processing and Management*, (42):387–406, 2006.

[5] H. Chen and S. T. Dumais. Bringing order to the web: automatically categorizing search results. In *Proc. of CHI-00, ACM International Conference on Human Factors in Computing Systems*, pages 145–152, 2000.

[6] P. J. Garcés, J. A. Olivas, and F. P. Romero. Concept-matching ir systems versus word-matching information retrieval systems: Considering fuzzy interrelations for indexing web pages. *JASIS&T*, 57(4):564–576, 2006.

[7] K. Golub. Automated subject classification of textual Web documents. *Journal of Documentation*, 62(3):350–371, 2006.

[8] K. Golub. Automated subject classification of textual web pages, based on a controlled vocabulary: challenges and recommendations. *New review of hypermedia and multimedia*, 12(1):11–27, June 2006. Special issue on knowledge organization systems and services.

[9] K. Golub. The role of different thesauri terms in automated subject classification of text. In *IEEE/WIC/ACM International Conference on Web Intelligence*, Dec. 2006.

[10] K. Golub and A. Ardö. Importance of HTML Structural Elements in Automated Subject Classification. In A. Rauber, S. Christodoulakis, and A. M. Tjoa, editors, *9th European Conference on Research and Advanced Technology for Digital Libraries - ECDL 2005*, volume 3652 of *Lecture Notes in Computer Science*, pages 368 – 378. Springer, Sept. 2005. Manuscript at: http://www.it.lth.se/knowlib/publ/ECDL2005.pdf.

[11] K. Golub, A. Ardö, D. Mladenic, and M. Grobelnik. Comparing and Combining Two Approaches to Automated Subject Classification of Text. In J. Gonzalo, C. Thanos, M. F. Verdejo, and R. C. Carrasco, editors, *10th European Conference on Research and Advanced Technology for Digital Libraries - ECDL 2006*, volume 4172 of *Lecture Notes in Computer Science*, pages 467–470. Springer, Sept. 2006.

[12] P. Ingwersen and K. Järvelin. *The turn: integration of information seeking and retrieval in context.* Springer, Dordrecht, The Netherlands, 2005.

[13] F. W. Lancaster. *Indexing and abstracting in theory and practice.* Facet, London, 2003. 3rd ed.

[14] D. D. Lewis, Y. Yang, T. Rose, and F. Li. Rcv1: A new benchmark collection for text categorization research. *The Journal of Machine Learning Research*, (5):361–397, 2004.

[15] O. Medelyan and I. Witten. Thesaurus based automatic keyphrase indexing. In *Proceedings of the Sixth ACM/IEEE Joint Conference on Digital Libraries, JCDL 06*, pages 296–297, 2006.

[16] H. A. Olson and J. J. Boll. *Subject analysis in online catalogs*. Englewood, CO: Libraries Unlimited, 2001. 2nd ed.

[17] F. Sebastiani. Machine learning in automated text categorization. *ACM Computing Surveys*, 34(1):1–47, 2002.

[18] E. Svenonius. *The intellectual foundations of information organization*. MIT Press, Cambridge, MA, USA, 2000.

[19] R. R. Trujilo. Simulation tool to study focused web crawling strategies. Master's thesis, Dept. of Information Technology, Lund University, P.O. Box 118, S-221 00 Lund, Sweden, Mar. 2006. http://combine.it.lth.se/CrawlSim/CrawlSim.pdf.

[20] Y. Yang. An evaluation of statistical approaches to text categorization. *Journal of Information Retrieval*, (1):67–88, 1999.

Part II
Gory details

8   Frequently asked questions

1. What does the message 'Wide character in subroutine entry ...' mean?

   That something is horribly wrong with the character encoding of this page.

2. What does the message 'Parsing of undecoded UTF-8 will give garbage when decoding entities ...' mean?

   That something is wrong with character decoding of this page.

3. I can't figure out how to restrict the crawler to pages below '`http://www.foo.com/bar/`'?

   Put an appropriate regular expression in the <allow> section of the configuration file. Appropriate means a Perl regular expression, which means that you have to escape special characters. Try with
   `URL http:\/\/www\.foo\.com\/bar\/`

4. I have a simple configuration variable set, but Combine does not obey it?

   Check that there are not 2 instances of the same simple configuration variable in the same configuration file. Unfortunately this will break configuration loading.

5. If there are multiple <allow> entries, must an URL fit all or any of them?

   A match to any of the entries will make that URL allowable for crawling. You can use any mix of HOST: and URL entries

6. It would also be nice to be able to crawl local files.

   Presently the crawler only accepts HTTP, HTTPS, and FTP as protocols.

7. Crawling of a single host is VERY slow. Is there some way for me to speed the crawler up?

   Yes it's one of the built-in limitations to keep the crawler beeing 'nice'. It will only access a particular server once every 60 seconds by default. You can change the default by adjusting the following configuration variables, but please keep in mind that you increase the load on the server.
   WaitIntervalSchedulerGetJcf=2
   WaitIntervalHost = 5

8. Is it possible to crawl only one single web-page?

   Use the command:
   `combine --jobname XXX --harvesturl http://www.foo.com/bar.html`

9. How can I crawl a fixed number of link steps from a set of seed pages? For example one web-page and all local links on that web-page (and not any further?

   Initialize the database and load the seed pages. Turn of automatic recycling of links by setting the simple configuration variable 'AutoRecycleLinks' to 0.

Start crawling and stop when 'combineCtrl -jobname XXX howmany' equals 0.

Handle recycling manually using 'combineCtrl, with action 'recyclelinks'. (Give the command `combineCtrl -jobname XXX recyclelinks`')

Iterate to the depth of your liking.

10. I run combineINIT but the configuration directory is not created?

You need to run combineINIT as root, due to file protection permissions.

11. Where are the logs?

They are stored in the SQL database <jobname> in the table `log`.

12. What are the main differences between Std (`classifyPlugIn = Combine::Check_record`) and PosCheck (`classifyPlugIn = Combine::PosCheck_record`) algorithms for automated subject classification?

Std can handle Perl regular expressions in terms and does not take into account if the term is found in the beginning or end of the document. PosCheck can't handle Perl regular expressions but is faster, and takes word position and proximity into account.

For detailed descriptions see sections Algorithm 1 (**??**) Algorithm 2 (**??**).

13. I don't understand what this means. Can you explain it to me ? Thank you !

```
40: sundew[^\s]*=CP.Drosera
40: tropical pitcher plant=CP.Nepenthes
```

It's part of the topic definition (term list) for the topic 'Carnivorous plants'. It's well described in the documentation, please see section **??**. The strange characters are Perl regular expressions mostly used for truncation etc.

14. I want to get all pages about "icecream" from "www.yahoo.com". And I don't have clear idea about how to write the topic definition file. Can you show me an example?

So for getting all pages about 'icecream' from 'www.yahoo.com' you have to:

(a) write a topic definition file according to the format above, eg containing topic specific terms. The file is essential a list of terms relevant for the topic. Format of the file is "numeric_importance: term=TopicClass" e.g. "100: icecream=YahooIce" (Say you call your topic 'YahooIce'). A few terms might be:

```
100: icecream=YahooIce
100: ice cone=YahooIce
```

and so on stored in a file called say TopicYahooIce.txt

(b) Initialization
`sudo combineINIT -jobname cptest -topic TopicYahooIce.txt`

(c) Edit the configuration to only allow crawling of www.yahoo.com Change the <allow> part in /etc/combine/focustest/combine.cfg from

```
#use either URL or HOST: (obs ':') to match regular expressions to either the
#full URL or the HOST part of a URL.
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: .*$
</allow>

to

#use either URL or HOST: (obs ':') to match regular expressions to either the
#full URL or the HOST part of a URL.
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: www\.yahoo\.com$
</allow>
```

(d) Load some good seed URLs

(e) Start 1 harvesting process

15. Why load some good seeds URLs and what the seeds URLs mean.

    This is just a way of telling the crawler where to start.

16. My problem is that the installation there requires root access, which I cannot get. Is there a way of running Combine without requiring any root access?

    The are three things that are problematic

    (a) Configurations are stored in /etc/combine/...

    (b) Runtime PID files are stored in /var/run/combine

    (c) You have to be able to create MySQL databases accessible by combine

    If you take the source and look how the tests (make test) are made you might find a way to fix the first. Though this probably involves modifying the source - maybe only the Combine/Config.pm

    The second is strictly not necessary and it will run even if /var/run /combine does not exist, although not the command `combineCtrl --jobname XXX kill`

    On the other hand the third is necessary and I can't think of a way around it except making a local installation of MySQL and use that.

17. What does the following entries from the log table mean?

    (a) `| 5409 | HARVPARS 1_zltest | 2006-07-14 15:08:52 | M500; SD empty, sleep 20 second..`
        This means that there are no URLs ready for crawling (SD empty). Also you can use combineCtrl to see current status of ready queue etc

    (b) `| 7352 | HARVPARS 1_wctest | 2006-07-14 17:00:59 | M500; urlid=1; netlocid=1; http:/`
        Crawler process 7352 got a URL (http://www.shanghaidaily.com/) to check (1_wctest is a just a name non significant) M500 is a sequence number for an individual crawler starting at 500 and when it reaches 0 this crawler process is killed and another is created. urlid and netlocid are internal identifiers used in the MySQL tables.

(c) | 7352 | HARVPARS 1_wctest | 2006-07-14 17:01:10 | M500; RobotRules OK, OK

Crawler process have checked that this URL (identified earlier in the log by pid=7352 and M500) can be crawled according to the Robot Exclusion protocol.

(d) | 7352 | HARVPARS 1_wctest | 2006-07-14 17:01:10 | M500; HTTP(200 = "OK") => OK

It has fetched the page (identified earlier in the log by pid=7352 and M500) OK

(e) | 7352 | HARVPARS 1_wctest | 2006-07-14 17:01:10 | M500; Doing: text/html;200;0F0610:

It is processing the page (in the format text/html) to see if it is of topical interest 0F061033DAF69587170F8E285E950120 is the MD5 checksum of the page

18. In fact, I want to know which crawled URLs are corresponding to the certain topic class such as CP.Aldrovanda . Can you tell me how can I know ?

You have to get into the raw MySQL database and perform a query like

SELECT urls.urlstr FROM urls,recordurl,topic WHERE urls.urlid=recordurl.urlid AND recordurl.recordid=topic.recordid AND topic.notation='CP.Aldrovanda';

Table urls contain all URLs seen by the crawler. Table recordurl connect urlid to recordid. recordid is used in all tables with data from the crawled Web pages.

19. What is the meaning of the item "ALL" in the notation column of the topic table?

If you use multiple topics in your topic-definition (ie the string after '=') then all the relevant topic scores for this page is summed and given the topic notation 'ALL'.

Just disregard it if you only use one topic-class.

20. Combine should crawl all pages underneath www.geocities.com/boulevard/newyork/, but not go outside the domain (i.e. going to www.yahoo.com) but also not going higher in position (i.e. www.geocities.com/boulevard/atlanta/).
Is it possible to set up Combine like this?

Yes, change the <allow>-part of your configuration file combine.cfg to select what URLs should be allowed for crawling (by default everything is allowed). See also section **??**.

So change

```
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: .*$
</allow>
```

to something like

```
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
URL http:\/\/www\.geocities\.com\/boulevard\/newyork\/
</allow>
```

(the backslashes are needed since these patterns are in fact Perl regular expressions)

9 Configuration variables

9.1 Name/value configuration variables

9.1.1 analysePlugin

Used by: utilPlugIn.pm.svn-base; utilPlugIn.pm

9.1.2 AutoRecycleLinks

Default value = 1

Description: Enable(1)/disable(0) automatic recycling of new links

Used by: SD_SQL.pm.svn-base; SD_SQL.pm

9.1.3 baseConfigDir

Default value = /etc/combine

Description: Base directory for configuration files; initialized by Config.pm

Used by: combineExport.svn-base; FromHTML.pm; combineExport; FromHTML.pm.svn-base

Set by: Config.pm; Config.pm.svn-base

9.1.4 classifyPlugIn

Default value = Combine::Check_record

Description: Which topic classification PlugIn module algorithm to use
Combine::Check_record and Combine::PosCheck_record included by default
NEW SVM classifier: Combine::classifySVM
see classifyPlugInTemplate.pm and documentation to write your own

Used by: combineReClassify.svn-base; combine; combineReClassify; combine.svn-base

9.1.5 configDir

Default value = NoDefaultValue

Description: Directory for job specific configuration files; taken from 'jobname'

Used by: classifySVM.pm.svn-base; combineUtil; utilPlugIn.pm.svn-base; classifySVM.pm; PosCheck_record.pm; Check_record.pm; Check_record.pm.svn-base; PosCheck_record.pm.svn-base; utilPlugIn.pm; combineUtil.svn-base

Set by: Config.pm; Config.pm.svn-base

9.1.6 doAnalyse

Default value = 1

Description: Enable(1)/disable(0) analysis of genre, language

Used by: combine; combine.svn-base

9.1.7 doCheckRecord

Description: Enable(1)/disable(0) topic classification (focused crawling)
Generated by combineINIT based on –topic parameter

Used by: combine; combine.svn-base; combineReClassify.svn-base; combineReClassify

9.1.8 doOAI

Default value = 1

Description: Use(1)/do not use(0) OAI record status keeping in SQL database

Used by: MySQLhdb.pm.svn-base; MySQLhdb.pm

9.1.9 extractLinksFromText

Default value = 1

Description: Extract(1)/do not extract(0) links from plain text

Used by: combine; combine.svn-base

9.1.10 HarvesterMaxMissions

Default value = 500

Description: Number of pages to process before restarting the harvester

Used by: combine; combine.svn-base

9.1.11 HarvestRetries

Default value = 5

Used by: combine; combine.svn-base

9.1.12 httpProxy

Default value = NoDefaultValue

Description: Use a proxy server if this is defined (default no proxy)

Used by: UA.pm; UA.pm.svn-base

9.1.13 LogHandle

Used by: classifySVM.pm.svn-base; FromHTML.pm; classifySVM.pm; PosCheck_record.pm; Check_record.pm; Check_record.pm.svn-base; PosCheck_record.pm.svn-base; FromHTML.pm.svn-base

Set by: combineReClassify.svn-base; combine; combineReClassify; combine.svn-base

9.1.14 Loglev

Description: Logging level (0 (least) - 10 (most))

Used by: combine; combine.svn-base

9.1.15 maxUrlLength

Default value = 250

Description: Maximum length of a URL; longer will be silently discarded

Used by: selurl.pm.svn-base; selurl.pm

9.1.16 MySQLdatabase

Default value = NoDefaultValue

Description: Identifies MySQL database name, user and host

Used by: Config.pm; Config.pm.svn-base

9.1.17 MySQLfulltext

Description: Enable(1)/disable(0) fulltext-index in MySQL table search

Used by: MySQLhdb.pm.svn-base; MySQLhdb.pm

9.1.18 MySQLhandle

Used by: combineSVM.svn-base; combineExport; MySQLhdb.pm.svn-base; combineReClassify.svn-base; combineUtil.svn-base; classifySVM.pm.svn-base; combineUtil; RobotRules.pm.svn-base; LogSQL.pm; combine; combineExport.svn-base; classifySVM.pm; RobotRules.pm; combine.svn-base; SD_SQL.pm; combineRank; combineRank.svn-base; LogSQL.pm.svn-base; XWI2XML.pm; combineSVM; XWI2XML.pm.svn-base; MySQLhdb.pm; SD_SQL.pm.svn-base; combineReClassify

Set by: Config.pm; Config.pm.svn-base

9.1.19 Operator-Email

Default value = "YourEmailAdress@YourDomain"

Description: Please change

Used by: RobotRules.pm.svn-base; RobotRules.pm; UA.pm; UA.pm.svn-base

9.1.20 Password

Default value = "XxXxyYzZ"

Description: Password not used yet. (Please change)

9.1.21 relTextPlugin

Used by: FromHTML.pm.svn-base; FromHTML.pm

9.1.22 saveHTML

Default value = 1

Description: Store(1)/do not store(0) the raw HTML in the database

Used by: MySQLhdb.pm.svn-base; MySQLhdb.pm

9.1.23 SdqRetries

Default value = 5

9.1.24 SummaryLength

Description: How long the summary should be. Use 0 to disable the summarization code

Used by: FromHTML.pm.svn-base; FromHTML.pm

9.1.25 SVMmodel

Default value = NoDefaultValue

Description: Filename for the SVM model

Used by: classifySVM.pm.svn-base; classifySVM.pm

9.1.26 UAtimeout

Default value = 30

Description: Time in seconds to wait for a server to respond

Used by: UA.pm; UA.pm.svn-base

9.1.27 UserAgentFollowRedirects

Description: User agent handles redirects (1) or treat redirects as new links (0)

Used by: UA.pm; UA.pm.svn-base

9.1.28 UserAgentGetIfModifiedSince

Default value = 1

Description: If we have seen this page before use Get-If-Modified (1) or not (0)

Used by: UA.pm; UA.pm.svn-base

9.1.29 useTidy

Description: Use(1)/do not use(0) Tidy to clean the HTML before parsing it

Used by: FromHTML.pm.svn-base; FromHTML.pm

9.1.30 WaitIntervalExpirationGuaranteed

Default value = 315360000

Used by: UA.pm; UA.pm.svn-base

9.1.31 WaitIntervalHarvesterLockNotFound

Default value = 2592000

Used by: combine; combine.svn-base

9.1.32 WaitIntervalHarvesterLockNotModified

Default value = 2592000

Used by: combine; combine.svn-base

9.1.33 WaitIntervalHarvesterLockRobotRules

Default value = 2592000

Used by: combine; combine.svn-base

9.1.34 WaitIntervalHarvesterLockSuccess

Default value = 1000000

Description: Time in seconds after succesfull download before allowing a page to be downloaded again (around 11 days)

Used by: combine; combine.svn-base

9.1.35 WaitIntervalHarvesterLockUnavailable

Default value = 86400

Used by: combine; combine.svn-base

9.1.36 WaitIntervalHost

Default value = 60

Description: Minimum time between accesses to the same host. Must be positive

Used by: SD_SQL.pm; SD_SQL.pm.svn-base

9.1.37 WaitIntervalRrdLockDefault

Default value = 86400

Used by: RobotRules.pm.svn-base; RobotRules.pm

9.1.38 WaitIntervalRrdLockNotFound

Default value = 345600

Used by: RobotRules.pm.svn-base; RobotRules.pm

9.1.39 WaitIntervalRrdLockSuccess

Default value = 345600

Used by: RobotRules.pm.svn-base; RobotRules.pm

9.1.40 WaitIntervalSchedulerGetJcf

Default value = 20

Description: Time in seconds to wait before making a new reschedule if a reschedule results in an empty ready que

Used by: combine; combine.svn-base

9.1.41 ZebraHost

Default value = NoDefaultValue

Description: Direct connection to Zebra indexing - for SearchEngine-in-a-box (default no connection)

Used by: MySQLhdb.pm.svn-base; MySQLhdb.pm

9.2 Complex configuration variables

9.2.1 allow

Description: use either URL or HOST: (obs ':') to match regular expressions to either the full URL or the HOST part of a URL.
Allow crawl of URLs or hostnames that matches these regular expressions

Used by: selurl.pm.svn-base; selurl.pm

9.2.2 binext

Description: Extensions of binary files

Used by: UA.pm; UA.pm.svn-base

9.2.3 converters

Description: Configure which converters can be used to produce a XWI object
Format:
1 line per entry
each entry consists of 3 ';' separated fields
Entries are processed in order and the first match is executed
external converters have to be found via PATH and executable to be considered a match
the external converter command should take a filename as parameter and convert that file
the result should be comming on STDOUT
mime-type ; External converter command ; Internal converter

Used by: UA.pm; combine; UA.pm.svn-base; combine.svn-base

9.2.4 exclude

Description: Exclude URLs or hostnames that matches these regular expressions
default: CGI and maps
default: binary files
default: Unparsable documents
default: images
default: other binary formats
more excludes in the file config_exclude (automatically updated by other programs)

Used by: selurl.pm.svn-base; selurl.pm

9.2.5 serveralias

Description: List of servernames that are aliases are in the file ./config_serveralias
(automatically updated by other programs)
use one server per line
example
www.100topwetland.com www.100wetland.com
means that www.100wetland.com is replaced by www.100topwetland.com during URL normalization

9.2.6 sessionids

Description: patterns to recognize and remove sessionids in URLs

9.2.7 url

Description: url is just a conatiner for all URL related configuration patterns

Used by: selurl.pm.svn-base; Config.pm; Config.pm.svn-base; selurl.pm

10 Module dependences

10.1 Programs

10.1.1 Check_record.pm.svn-base

Uses: Combine::XWI; Combine::LoadTermList; Combine::utilPlugIn; Combine::Matcher; Combine::Config;

Used by:

10.1.2 CleanXML2CanDoc.pm.svn-base

Uses: Alvis::Canonical;

Used by:

10.1.3 Config.pm.svn-base

Uses: Config::General; DBI;

Used by:

10.1.4 DataBase.pm.svn-base

Uses: Combine::MySQLhdb; Digest::MD5; Encode; Combine::selurl;

Used by:

10.1.5 FromHTML.pm.svn-base

Uses: Combine::Config; HTTP::Date; URI; URI::Escape; HTML::Entities; Encode; HTML::Tidy; Combine::HTMLExtractor;

Used by:

10.1.6 FromImage.pm.svn-base

Uses: Image::ExifTool;

Used by:

10.1.7 HTMLExtractor.pm.svn-base

Uses: HTML::TokeParser; URI; Data::Dumper;

Used by:

10.1.8 LoadTermList.pm.svn-base

Uses: DBI; Lingua::Stem;

Used by:

10.1.9 LogSQL.pm.svn-base

Uses: Combine::Config;

Used by:

10.1.10 Matcher.pm.svn-base

Uses: HTML::Entities;

Used by:

10.1.11 MySQLhdb.pm.svn-base

Uses: Combine::XWI; HTTP::Date; Encode; Combine::Config; Combine::selurl; Combine::Zebra; Combine::Zebra;

Used by:

10.1.12 PosCheck_record.pm.svn-base

Uses: Combine::LoadTermList; Combine::utilPlugIn; Combine::PosMatcher; Combine::Config;

Used by:

10.1.13 PosMatcher.pm.svn-base

Uses: Combine::XWI2XML; HTML::Entities;

Used by:

10.1.14 RobotRules.pm.svn-base

Uses: Combine::Config; Combine::UA;

Used by:

10.1.15 SD_SQL.pm.svn-base

Uses: Combine::Config; Combine::selurl; DBI;

Used by:

10.1.16 UA.pm.svn-base

Uses: Combine::Config; LWP::UserAgent; HTTP::Date;

Used by:

10.1.17 XWI.pm.svn-base

Uses: HTML::Entities; Combine::XWI;

Used by:

10.1.18   XWI2XML.pm.svn-base

Uses: Combine::XWI; Encode; Combine::Config; Compress::Zlib; MIME::Base64; Combine::CleanXML2CanDoc;

Used by:

10.1.19   Zebra.pm.svn-base

Uses: Combine::XWI2XML; ZOOM;

Used by:

10.1.20   classifySVM.pm.svn-base

Uses: Combine::XWI; Combine::Config; Combine::MySQLhdb; Combine::utilPlugIn;

Used by:

10.1.21   combine

Uses: Combine::Config; Combine::XWI; Combine::UA; Combine::RobotRules; Combine::LogSQL; Combine::FromHTML; Combine::FromImage; Combine::FromTeX; Combine::utilPlugIn; Combine::DataBase; HTTP::Date; HTTP::Status; URI::URL; Getopt::Long; Combine::SD_SQL;

10.1.22   combine.svn-base

Uses: Combine::Config; Combine::XWI; Combine::UA; Combine::RobotRules; Combine::LogSQL; Combine::FromHTML; Combine::FromImage; Combine::FromTeX; Combine::utilPlugIn; Combine::DataBase; HTTP::Date; HTTP::Status; URI::URL; Getopt::Long; Combine::SD_SQL;

10.1.23   combineCtrl

Uses: Getopt::Long; Combine::SD_SQL; Combine::Config;

10.1.24   combineCtrl.svn-base

Uses: Getopt::Long; Combine::SD_SQL; Combine::Config;

10.1.25   combineExport

Uses: Combine::MySQLhdb; Combine::Config; Combine::XWI2XML; DBI; HTTP::Date; Encode; Getopt::Long; Alvis::Pipeline; XML::LibXSLT; XML::LibXML;

10.1.26   combineExport.svn-base

Uses: Combine::MySQLhdb; Combine::Config; Combine::XWI2XML; DBI; HTTP::Date; Encode; Getopt::Long; Alvis::Pipeline; XML::LibXSLT; XML::LibXML;

10.1.27   combineINIT

Uses: Getopt::Long; Combine::Config; DBI; HTML::Tidy;

10.1.28   combineINIT.svn-base

Uses: Getopt::Long; Combine::Config; DBI; HTML::Tidy;

10.1.29   combineRank

Uses: Getopt::Long; Combine::Config; Combine::GraphAlgorithm; DBI;

10.1.30   combineRank.svn-base

Uses: Getopt::Long; Combine::Config; Combine::GraphAlgorithm; DBI;

10.1.31   combineReClassify

Uses: Combine::Config; Combine::LogSQL; Combine::DataBase; Combine::MySQLhdb; Combine::XWI; DBI; Getopt::Long;

10.1.32   combineReClassify.svn-base

Uses: Combine::Config; Combine::LogSQL; Combine::DataBase; Combine::MySQLhdb; Combine::XWI; DBI; Getopt::Long;

10.1.33   combineSVM

Uses: warnings; Combine::Config; Combine::selurl; Combine::MySQLhdb; Combine::utilPlugIn; Getopt::Long; Algorithm::SVMLight;

10.1.34   combineSVM.svn-base

Uses: warnings; Combine::Config; Combine::selurl; Combine::MySQLhdb; Combine::utilPlugIn; Getopt::Long; Algorithm::SVMLight;

10.1.35   combineUtil

Uses: Getopt::Long; Combine::Config; Combine::SD_SQL; Combine::MySQLhdb; Combine::MySQLhdb; Net::hostent;

10.1.36   combineUtil.svn-base

Uses: Getopt::Long; Combine::Config; Combine::SD_SQL; Combine::MySQLhdb; Combine::MySQLhdb; Net::hostent;

10.1.37   selurl.pm.svn-base

Uses: URI; Combine::Config;

Used by:

10.1.38   utilPlugIn.pm.svn-base

Uses: Combine::XWI; Combine::MySQLhdb; Combine::LoadTermList; Geo::IP; Locale::Country; Lingua::Identify; Algorithm::SVMLight;

Used by:

10.2   Library modules

10.2.1   Check_record.pm

Uses: Combine::XWI; Combine::LoadTermList; Combine::utilPlugIn; Combine::Matcher; Combine::Config;

Used by:

10.2.2   CleanXML2CanDoc.pm

Uses: Alvis::Canonical;

Used by: Combine::XWI2XML; Combine::.svn::text-base::XWI2XML.pm.svn-base;

10.2.3   Config.pm

Uses: Config::General; DBI;

Used by: combineExport; combineINIT; combineCtrl; combineReClassify; combineSVM; combineRank; combineUtil; combine; combineINIT.svn-base; combineCtrl.svn-base; combineReClassify.svn-base; combineSVM.svn-base; combineRank.svn-base; combineUtil.svn-base; combine.svn-base; combineExport.svn-base; Combine::Check_record; Combine::UA; Combine::selurl; Combine::XWI2XML; Combine::MySQLhdb; Combine::SD_SQL; Combine::FromHTML; Combine::PosCheck_record; Combine::RobotRules; Combine::LogSQL; Combine::classifySVM; Combine::.svn::text-base::UA.pm.svn-base; Combine::.svn::text-base::selurl.pm.svn-base; Combine::.svn::text-base::XWI2XML.pm.svn-base; Combine::.svn::text-base::Check_record.pm.svn-base; Combine::.svn::text-base::MySQLhdb.pm.svn-base; Combine::.svn::text-base::SD_SQL.pm.svn-base; Combine::.svn::text-base::FromHTML.pm.svn-base; Combine::.svn::text-base::PosCheck_record.pm.svn-base; Combine::.svn::text-base::RobotRules.pm.svn-base; Combine::.svn::text-base::LogSQL.pm.svn-base; Combine::.svn::text-base::classifySVM.pm.svn-base;

10.2.4   DataBase.pm

Uses: Combine::MySQLhdb; Digest::MD5; Encode; Combine::selurl;

Used by: combineReClassify; combine; combineReClassify.svn-base; combine.svn-base;

10.2.5   FromHTML.pm

Uses: Combine::Config; HTTP::Date; URI; URI::Escape; HTML::Entities; Encode; HTML::Tidy; Combine::HTMLExtractor;

Used by: combine; combine.svn-base;

10.2.6 FromImage.pm

Uses: Image::ExifTool;

Used by: combine; combine.svn-base;

10.2.7 HTMLExtractor.pm

Uses: HTML::TokeParser; URI; Data::Dumper;

Used by: Combine::FromHTML; Combine::.svn::text-base::FromHTML.pm.svn-base;

10.2.8 LoadTermList.pm

Uses: DBI; Lingua::Stem;

Used by: Combine::Check_record; Combine::utilPlugIn; Combine::PosCheck_record; Combine::.svn::text-base::Check_record.pm.svn-base; Combine::.svn::text-base::utilPlugIn.pm.svn-base; Combine::.svn::text-base::PosCheck_record.pm.svn-base;

10.2.9 LogSQL.pm

Uses: Combine::Config;

Used by: combineReClassify; combine; combineReClassify.svn-base; combine.svn-base;

10.2.10 Matcher.pm

Uses: HTML::Entities;

Used by: Combine::Check_record; Combine::.svn::text-base::Check_record.pm.svn-base;

10.2.11 MySQLhdb.pm

Uses: Combine::XWI; HTTP::Date; Encode; Combine::Config; Combine::selurl; Combine::Zebra; Combine::Zebra;

Used by: combineExport; combineReClassify; combineSVM; combineUtil; combineUtil; combineReClassify.svn-base; combineSVM.svn-base; combineUtil.svn-base; combineUtil.svn-base; combineExport.svn-base; Combine::DataBase; Combine::utilPlugIn; Combine::classifySVM; Combine::.svn::text-base::DataBase.pm.svn-base; Combine::.svn::text-base::utilPlugIn.pm.svn-base; Combine::.svn::text-base::classifySVM.pm.svn-base;

10.2.12 PosCheck_record.pm

Uses: Combine::LoadTermList; Combine::utilPlugIn; Combine::PosMatcher; Combine::Config;

Used by:

10.2.13 PosMatcher.pm

Uses: Combine::XWI2XML; HTML::Entities;

Used by: Combine::PosCheck_record; Combine::.svn::text-base::PosCheck_record.pm.svn-base;

10.2.14 RobotRules.pm

Uses: Combine::Config; Combine::UA;

Used by: combine; combine.svn-base;

10.2.15 SD_SQL.pm

Uses: Combine::Config; Combine::selurl; DBI;

Used by: combineCtrl; combineUtil; combine; combineCtrl.svn-base; combineUtil.svn-base; combine.svn-base;

10.2.16 UA.pm

Uses: Combine::Config; LWP::UserAgent; HTTP::Date;

Used by: combine; combine.svn-base; Combine::RobotRules; Combine::.svn::text-base::RobotRules.pm.svn-base;

10.2.17 XWI.pm

Uses: HTML::Entities; Combine::XWI;

Used by: combineReClassify; combine; combineReClassify.svn-base; combine.svn-base; Combine::Check_record; Combine::XWI; Combine::XWI2XML; Combine::MySQLhdb; Combine::utilPlugIn; Combine::classifySVM; Combine::.svn::text-base::XWI.pm.svn-base; Combine::.svn::text-base::XWI2XML.pm.svn-base; Combine::.svn::text-base::Check_record.pm.svn-base; Combine::.svn::text-base::MySQLhdb.pm.svn-base; Combine::.svn::text-base::utilPlugIn.pm.svn-base; Combine::.svn::text-base::classifySVM.pm.svn-base;

10.2.18 XWI2XML.pm

Uses: Combine::XWI; Encode; Combine::Config; Compress::Zlib; MIME::Base64; Combine::CleanXML2CanDoc;

Used by: combineExport; combineExport.svn-base; Combine::PosMatcher; Combine::Zebra; Combine::.svn::text-base::PosMatcher.pm.svn-base; Combine::.svn::text-base::Zebra.pm.svn-base;

10.2.19 Zebra.pm

Uses: Combine::XWI2XML; ZOOM;

Used by: Combine::MySQLhdb; Combine::MySQLhdb; Combine::.svn::text-base::MySQLhdb.pm.svn-base; Combine::.svn::text-base::MySQLhdb.pm.svn-base;

10.2.20 classifySVM.pm

Uses: Combine::XWI; Combine::Config; Combine::MySQLhdb; Combine::utilPlugIn;

Used by:

10.2.21   selurl.pm

Uses: URI; Combine::Config;

Used by: combineSVM; combineSVM.svn-base; Combine::DataBase; Combine::MySQLhdb; Combine::SD_SQL; Combine::.svn::text-base::DataBase.pm.svn-base; Combine::.svn::text-base::MySQLhdb.pm.svn-base; Combine::.svn::text-base::SD_SQL.pm.svn-base;

10.2.22   utilPlugIn.pm

Uses: Combine::XWI; Combine::MySQLhdb; Combine::LoadTermList; Geo::IP; Locale::Country; Lingua::Identify; Algorithm::SVMLight;

Used by: combineSVM; combine; combineSVM.svn-base; combine.svn-base; Combine::Check_record; Combine::PosCheck_record; Combine::classifySVM; Combine::.svn::text-base::Check_record.pm.svn-base; Combine::.svn::text-base::PosCheck_record.pm.svn-base; Combine::.svn::text-base::classifySVM.pm.svn-base;

10.3   External modules

These are the (non base) Perl modules Combine depend on. The modules marked with a '*' are not critical.

```
Algorithm::SVMLight
Algorithm::SVMLight *
Alvis::Canonical
Alvis::Pipeline *
Compress::Zlib
Config::General
DBI
Data::Dumper *
Digest::MD5
Encode
Geo::IP
Getopt::Long
HTML::Entities
HTML::Tidy *
HTML::TokeParser
HTTP::Date
HTTP::Status
Image::ExifTool
LWP::UserAgent
Lingua::Identify
Lingua::Stem
Locale::Country
MIME::Base64
Net::hostent
URI
URI::Escape
URI::URL
```

```
XML::LibXML
XML::LibXSLT
ZOOM *
warnings
```

Part III

A APPENDIX

A.1 Simple installation test

The following simple script is available in the `doc/InstallationTest.pl` file. It must be run as 'root' and tests that basic functions of the Combine installation works.

Basicly it creates and initializes a new jobname, crawls one specific test page and exports it as XML. This XML is then compared to a correct XML-record for that page.

A.1.1 InstallationTest.pl

```
use strict;
if ( $> != 0 ) {
    die("You have to run this test as root");
}

my $orec='';
while (<DATA>) { chop; $orec .= $_; }

$orec =~ s|<checkedDate>.*</checkedDate>||;
$orec =~ tr/\n\t //d;

my $olen=length($orec);
my $onodes=0;
while ( $orec =~ m/</g ) { $onodes++; }
print "ORIG Nodes=$onodes; Len=$olen\n";

our $jobname;
require './t/defs.pm';

system("combineINIT --jobname $jobname --topic /etc/combine/Topic_carnivor.txt > /dev/null");

system("combine --jobname $jobname --harvest http://combine.it.lth.se/CombineTests/Installati
open(REC,"combineExport --jobname $jobname |");
my $rec='';
while (<REC>) { chop; $rec .= $_; }
close(REC);
$rec =~ s|<checkedDate>.*</checkedDate>||;
$rec =~ tr/\n\t //d;

my $len=length($rec);
my $nodes=0;
while ( $rec =~ m/</g ) { $nodes++; }
print "NEW Nodes=$nodes; Len=$len\n";

my $OK=0;
```

58

```perl
if ($onodes == $nodes) { print "Number of XML nodes match\n"; }
else { print "Number of XML nodes does NOT match\n"; $OK=1; }
if ($olen == $len) {
  print "Size of XML match\n";
} else {
  $orec =~  s|<originalDocument.*</originalDocument>||s;
  $rec =~  s|<originalDocument.*</originalDocument>||s;
  if (length($orec) == length($rec)) { print "Size of XML match (after removal of 'originalDo
  else { print "Size of XML does NOT match\n"; $OK=1; }
}


if (($OK == 0) && ($orec eq $rec)) { print "All tests OK\n"; }
else { print "There might be some problem with your Combine Installation\n"; }


__END__
```

```xml
<?xml version="1.0" encoding="UTF-8"?>
<documentCollection version="1.1" xmlns="http://alvis.info/enriched/">
<documentRecord id="80AC707F96BC57DFEF78C815F6FABD57">
<acquisition>
<acquisitionData>
<modifiedDate>2006-12-05 13:20:25</modifiedDate>
<checkedDate>2006-10-03 9:06:42</checkedDate>
<httpServer>Apache/1.3.29 (Debian GNU/Linux) PHP/4.3.3</httpServer>
<urls>
    <url>http://combine.it.lth.se/CombineTests/InstallationTest.html</url>
  </urls>
</acquisitionData>
<originalDocument mimeType="text/html" compression="gzip" encoding="base64" charSet="UTF-8">
H4sIAAAAAAAA4WQsU7DMBCG9zzF4bmpBV2QcDKQVKJSKR2CEKObXBSrjm3sSSyFvT0yCQGJgusG/
//u+E1flU1G9HrfwUD3u4fh8v98VwFLOXzYF52VVzg+b9Q3n2wPLE9FRr+NA2UyDFGnMdyaQ1FqS
sgYIAOFrPRS2PymDgs+hRPRIEozsMWNnHN+tbwKD2hpCQxkrpDfqYrOdAjgtDYUVlN4G9HIFB3RT
qMPAvns6Ipfi26Au09e5I61Gh78aCT+IR947qDvpA1I2UJvexg6+CJxsMOad6/8kpkQiXB5XSWUC
BNsj/GGG4LBWrarhSw+OOiOIidZjmzGPeh15WL6ICS7zFUjT/AiuBXeRbwHj870/AeRYaTupAQAA
</originalDocument>
```
```xml
<canonicalDocument>
  <section>
    <section title="Installation test for Combine">
      <section>Installation test for Combine</section>
      <section>Contains some Carnivorous plant specific words like <ulink url="rel.html">Dros
<metaData>
    <meta name="title">Installation test for Combine</meta>
    <meta name="dc:format">text/html</meta>
    <meta name="dc:format">text/html; charset=iso-8859-1</meta>
    <meta name="dc:subject">Carnivorous plants</meta>
    <meta name="dc:subject">Drosera</meta>
    <meta name="dc:subject">Nepenthes</meta>
  </metaData>
<links>
```

```
    <outlinks>
      <link type="a">
        <anchorText>Drosera</anchorText>
        <location>http://combine.it.lth.se/CombineTests/rel.html</location>
      </link>
    </outlinks>
  </links>
<analysis>
<property name="topLevelDomain">se</property>
<property name="univ">1</property>
<property name="language">en</property>
<topic absoluteScore="1000" relativeScore="110526">
    <class>ALL</class>
  </topic>
<topic absoluteScore="375" relativeScore="41447">
    <class>CP.Drosera</class>
    <terms>drosera</terms>
  </topic>
<topic absoluteScore="375" relativeScore="41447">
    <class>CP.Nepenthes</class>
    <terms>nepenthe</terms>
  </topic>
<topic absoluteScore="250" relativeScore="27632">
    <class>CP</class>
    <terms>carnivorous plant</terms>
    <terms>carnivor</terms>
  </topic>
</analysis>
</acquisition>
</documentRecord>

</documentCollection>
```

A.2 Example topic filter plug in

This example gives more details on how to write a topic filter Plug-In.

A.2.1 classifyPlugInTemplate.pm

```
#Template for writing a classify PlugIn for Combine
#See documentation at http://combine.it.lth.se/documentation/

package classifyPlugInTemplate; #Change to your own module name

use Combine::XWI; #Mandatory
use Combine::Config; #Optional if you want to use the Combine configuration system

#API:
#  a subroutine named 'classify' taking a XWI-object as in parameter
```

```perl
#     return values: 0/1
#          0: record fails to meet the classification criteria, ie ignore this record
#          1: record is OK and should be stored in the database, and links followed by the crawl
sub classify {
  my ($self,$xwi) = @_;

  #utility routines to extract information from the XWI-object
  #URL (can be several):
   # $xwi->url_rewind;
   # my $url_str="";
   # my $t;
   # while ($t = $xwi->url_get) { $url_str .= $t . ", "; }

  #Metadata:
   #  $xwi->meta_rewind;
   #  my ($name,$content);
   #  while (1) {
   #     ($name,$content) = $xwi->meta_get;
   #     last unless $name;
   #     next if ($name eq 'Rsummary');
   #     next if ($name =~ /^autoclass/);
   #     $meta .= $content . " ";
   #  }

  #Title:
   #  $title = $xwi->title;

  #Headings:
   #  $xwi->heading_rewind;
   #  my $this;
   #  while (1) {
   #     $this = $xwi->heading_get or last;
   #     $head .= $this . " ";
   #  }

  #Text:
   #  $this = $xwi->text;
   #  if ($this) {
   #     $text = $$this;
   #  }

#############################
#Apply your classification algorithm here
#  assign $result a value (0/1)
#############################

  #utility routines for saving detailed results (optional) in the database. These data may app
  # in exported XML-records
```

```perl
  #Topic takes 5 parameters
  # $xwi->topic_add(topic_class_notation, topic_absolute_score, topic_normalized_score, topic_
  #  topic_class_notation, topic_terms, and algorithm_id are strings
  #    max length topic_class_notation: 50, algorithm_id: 25
  #  topic_absolute_score, and topic_normalized_score are integers
  #  topic_normalized_score and topic_terms are optional and may be replaced with 0, '' respe

  #Analysis takes 2 parameters
  # $xwi->robot_add(name,value);
  # both are strings with max length name: 15, value: 20

    # return true (1) if you want to keep the record
    # otherwise return false (0)

  return $result;
}

1;
```

A.3 Default configuration files

A.3.1 Global

```
#@#Default configuration values Combine system

#Direct connection to Zebra indexing - for SearchEngine-in-a-box (default no connection)
#@#ZebraHost = NoDefaultValue
ZebraHost =

#Enable(1)/disable(0) fulltext-index in MySQL table search
MySQLfulltext = 0

#Use a proxy server if this is defined (default no proxy)
#@#httpProxy = NoDefaultValue
httpProxy =

#Enable(1)/disable(0) automatic recycling of new links
AutoRecycleLinks = 1

#User agent handles redirects (1) or treat redirects as new links (0)
UserAgentFollowRedirects = 0

#Number of pages to process before restarting the harvester
HarvesterMaxMissions = 500

#Logging level (0 (least) - 10 (most))
Loglev = 0
```

```
#Enable(1)/disable(0) analysis of genre, language
doAnalyse = 1
analysePlugin =
relTextPlugin =

#How long the summary should be. Use 0 to disable the summarization code
SummaryLength   = 0

#Store(1)/do not store(0) the raw HTML in the database
saveHTML = 1

#Use(1)/do not use(0) Tidy to clean the HTML before parsing it
useTidy = 0

#Use(1)/do not use(0) OAI record status keeping in SQL database
doOAI = 1

#Extract(1)/do not extract(0) links from plain text
extractLinksFromText = 1

#Enable(1)/disable(0) topic classification (focused crawling)
#Generated by combineINIT based on --topic parameter
doCheckRecord = 0

#Which topic classification PlugIn module algorithm to use
#Combine::Check_record and Combine::PosCheck_record included by default
#NEW SVM classifier: Combine::classifySVM
#see classifyPlugInTemplate.pm and documentation to write your own
classifyPlugIn = Combine::Check_record

#Filename for the SVM model
#@#SVMmodel = NoDefaultValue
SVMmodel =

###Parameters for Std topic classification algorithm
###StdTitleWeight = 10 #
###StdMetaWeight = 4 #
###StdHeadingsWeight = 2 #
###StdCutoffRel = 10 #Class score must be above this % to be counted
###StdCutoffNorm = 0.2 #normalised cutoff for summed normalised score
###StdCutoffTot = 90 #non normalised cutoff for summed total score

###Parameters for Pos topic classification algorithm
###PosCutoffRel = 1 #Class score must be above this % to be counted
###PosCutoffNorm = 0.002 #normalised cutoff for summed normalised score
###PosCutoffTot = 1 #non normalised cutoff for summed total score

HarvestRetries                  = 5
```

```
SdqRetries                        = 5

#Maximum length of a URL; longer will be silently discarded
maxUrlLength = 250

#Time in seconds to wait for a server to respond
UAtimeout = 30

#If we have seen this page before use Get-If-Modified (1) or not (0)
UserAgentGetIfModifiedSince = 1

WaitIntervalExpirationGuaranteed = 315360000
WaitIntervalHarvesterLockNotFound = 2592000
WaitIntervalHarvesterLockNotModified = 2592000
WaitIntervalHarvesterLockRobotRules = 2592000
WaitIntervalHarvesterLockUnavailable = 86400
WaitIntervalRrdLockDefault = 86400
WaitIntervalRrdLockNotFound = 345600
WaitIntervalRrdLockSuccess = 345600

#Time in seconds after succesfull download before allowing a page to be downloaded again (aro
WaitIntervalHarvesterLockSuccess = 1000000

#Time in seconds to wait before making a new reschedule if a reschedule results in an empty re
WaitIntervalSchedulerGetJcf = 20

#Minimum time between accesses to the same host. Must be positive
WaitIntervalHost = 60

#Identifies MySQL database name, user and host
MySQLdatabase   = NoDefaultValue

#Base directory for configuration files; initialized by Config.pm
#@#baseConfigDir = /etc/combine

#Directory for job specific configuration files; taken from 'jobname'
#@#configDir = NoDefaultValue

<binext>
#Extensions of binary files
arff
au
avi
class
exe
fig
gif
gz
```

```
hqx
ica
jpeg
jpg
mat
mdb
mov
mp3
mpeg
mpg
msi
pcx
pdb
psd
ram
rar
raw
rmd
rmx
sav
sdd
shar
tar
tga
tgz
tif
tiff
vo
wav
wmv
wmz
xbm
xpm
z
zip
</binext>

<converters>
#Configure which converters can be used to produce a XWI object
#Format:
#  1 line per entry
#  each entry consists of 3 ';' separated fields
#
#Entries are processed in order and the first match is executed
#  external converters have to be found via PATH and executable to be considered a match
#  the external converter command should take a filename as parameter and convert that file
#   the result should be comming on STDOUT
#
```

```
# mime-type    ;    External converter command ; Internal converter

text/html ; ; GuessHTML
#Check this
www/unknown ; ; GuessHTML
text/plain ; ; GuessText
text/x-tex ;  tth -g -w1 -r <  ; TeXHTML
application/x-tex ;  tth -g -w1 -r < ; TeXHTML
text/x-tex ; untex -a -e -giso ; TeXText
application/x-tex ; untex -a -e -giso ; TeXText
text/x-tex ;   ; TeX
application/x-tex ; ; TeX
application/pdf ; pdftohtml -i -noframes -nomerge -stdout ; HTML
application/pdf ; pstotext ; Text
application/postscript ; pstotext ; Text
application/msword ; antiword -t ; Text
application/vnd.ms-excel ; xlhtml -fw ; HTML
application/vnd.ms-powerpoint ; ppthtml ; HTML
application/rtf ; unrtf --nopict --html ; HTML
image/gif ; ; Image
image/jpeg ; ; Image
image/tiff ; ; Image
</converters>

<url>
  <exclude>
    #Exclude URLs or hostnames that matches these regular expressions
    #Malformed hostnames
    HOST: http:\/\/\.
    HOST: \@
  </exclude>
</url>
```

A.3.2   Job specific

```
#Please change
Operator-Email      = "YourEmailAdress@YourDomain"

#Password not used yet. (Please change)
Password    = "XxXxyYzZ"

<converters>
#Configure which converters can be used to produce a XWI object
#Format:
#  1 line per entry
#  each entry consists of 3 ';' separated fields
#
#Entries are processed in order and the first match is executed
```

```
#   external converters have to be found via PATH and executable to be considered a match
#   the external converter command should take a filename as parameter and convert that file
#    the result should be comming on STDOUT
#
# mime-type    ;    External converter command ; Internal converter

application/pdf ; MYpdftohtml -i -noframes -nomerge -stdout ; HTML
</converters>


<url>
#List of servernames that are aliases are in the file ./config_serveralias
#    (automatically updated by other programs)
#use one server per line
#example
#www.100topwetland.com  www.100wetland.com
#  means that www.100wetland.com is replaced by www.100topwetland.com during URL normalizatio
<serveralias>
<<include config_serveralias>>
</serveralias>


#use either URL or HOST: (obs ':') to match regular expressions to
# either the full URL or the HOST part of a URL.
<allow>
#Allow crawl of URLs or hostnames that matches these regular expressions
HOST: .*$
</allow>


<exclude>
#Exclude URLs or hostnames that matches these regular expressions
# default: CGI and maps
URL cgi-bin|htbin|cgi|\?|\.map$|_vti_

# default: binary files
URL \.exe$|\.zip$|\.tar$|\.tgz$|\.gz$|\.hqx$|\.sdd$|\.mat$|\.raw$
URL \.EXE$|\.ZIP$|\.TAR$|\.TGZ$|\.GZ$|\.HQX$|\.SDD$|\.MAT$|\.RAW$

# default: Unparsable documents
URL \.shar$|\.rmx$|\.rmd$|\.mdb$|\.sav$
URL \.SHAR$|\.RMX$|\.RMD$|\.MDB$|\.SAV$

# default: images
URL \.gif$|\.jpg$|\.jpeg$|\.xpm$|\.tif$|\.tiff$|\.mpg$|\.mpeg$|\.mov$|\.wav$|\.au$|\.pcx$|\.x
URL \.GIF$|\.JPG$|\.JPEG$|\.XPM$|\.TIF$|\.TIFF$|\.MPG$|\.MPEG$|\.MOV$|\.WAV$|\.AU$|\.PCX$|\.X

# default: other binary formats
URL \.pdb$|\.class$|\.ica$|\.ram$|\.wmz$|\.arff$|\.rar$|\.vo$|\.fig$|\.mp3$|\.wmv$|\.avi$|\.m
URL \.PDB$|\.CLASS$|\.ICA$|\.RAM$|\.WMZ$|\.ARFF$|\.RAR$|\.VO$|\.FIG$|\.MP3$|\.WMV$|\.AVI$|\.M
```

```
#more excludes in the file config_exclude (automatically updated by other programs)
<<include config_exclude>>
</exclude>
<sessionids>
#patterns to recognize and remove sessionids in URLs
sessionid
lsessionid
jsessionid
SID
PHPSESSID
SessionID
BV_SessionID
</sessionids>
#url is just a conatiner for all URL related configuration patterns
</url>
```

A.4   SQL database

A.4.1   Create database

```
DROP DATABASE IF EXISTS $database;
CREATE DATABASE $database DEFAULT CHARACTER SET utf8;
USE $database;
```


A.4.2   Creating MySQL tables

```
All tables use UTF-8

Summary tables '^'=primary key, '*'=key:
TABLE hdb: recordid^, type, dates, server, title, ip, ...
TABLE links: recordid*, mynetlocid*, urlid*, netlocid*, linktype, anchor  (netlocid for urlid
TABLE meta: recordid*,  name, value
TABLE html: recordid^, html
TABLE analys: recordid*, name, value
TABLE topic: recordid*, notation*, absscore, relscore, terms, algorithm
TABLE localtags: netlocid, urlid, name, value
TABLE search: recordid^, stext*

(TABLE netlocalias: netlocid*, netlocstr^)
(TABLE urlalias: urlid*, urlstr^)
TABLE topichierarchy: node^, father*, notation*, caption, level
TABLE netlocs: netlocid^, netlocstr^, retries
TABLE urls: netlocid*, urlid^, urlstr^, path
TABLE urldb: netlocid*, urlid^, urllock, harvest*, retries, netloclock
TABLE newlinks urlid^, netlocid
TABLE recordurl: recordid*, urlid^, lastchecked, md5*, fingerprint*^
TABLE admin: status, queid, schedulealgorithm
TABLE log: pid, id, date, message
```

```
TABLE que: queid^, urlid, netlocid
TABLE robotrules: netlocid*, rule, expire
TABLE oai: recordid, md5^, date*, status
TABLE exports: host, port, last
```

A.4.3 Data tables

```
CREATE TABLE hdb (
  recordid int(11) NOT NULL default '0',
  type varchar(50) default NULL,
  title text,
  mdate timestamp NOT NULL,
  expiredate datetime default NULL,
  length int(11) default NULL,
  server varchar(50) default NULL,
  etag varchar(25) default NULL,
  nheadings int(11) default NULL,
  nlinks int(11) default NULL,
  headings mediumtext,
  ip mediumblob,
  PRIMARY KEY  (recordid)
) ENGINE=MyISAM AVG_ROW_LENGTH = 20000 MAX_ROWS = 10000000 DEFAULT CHARACTER SET=utf8;

CREATE TABLE html (
  recordid int(11) NOT NULL default '0',
  html mediumblob,
  PRIMARY KEY  (recordid)
) ENGINE=MyISAM AVG_ROW_LENGTH = 20000 MAX_ROWS = 10000000 DEFAULT CHARACTER SET=utf8;

CREATE TABLE links (
  recordid int(11) NOT NULL default '0',
  mynetlocid int(11) default NULL,
  urlid int(11) default NULL,
  netlocid int(11) default NULL,
  anchor text,
  linktype varchar(50) default NULL,
  KEY recordid (recordid),
  KEY urlid (urlid),
  KEY mynetlocid (mynetlocid),
  KEY netlocid (netlocid)
) ENGINE=MyISAM MAX_ROWS = 1000000000 DEFAULT CHARACTER SET=utf8;

CREATE TABLE meta (
  recordid int(11) NOT NULL default '0',
  name varchar(50) default NULL,
  value text,
  KEY recordid (recordid)
) ENGINE=MyISAM MAX_ROWS = 1000000000 DEFAULT CHARACTER SET=utf8;
```

```
CREATE TABLE analys (
  recordid int(11) NOT NULL default '0',
  name varchar(15) NOT NULL,
  value varchar(20),
  KEY recordid (recordid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE topic (
  recordid int(11) NOT NULL default '0',
  notation varchar(50) default NULL,
  abscore int(11) default NULL,
  relscore int(11) default NULL,
  terms text default NULL,
  algorithm varchar(25),
  KEY notation (notation),
  KEY recordid (recordid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE localtags (
  netlocid int(11) NOT NULL DEFAULT '0',
  urlid int(11) NOT NULL DEFAULT '0',
  name varchar(100) NOT NULL,
  value varchar(100) NOT NULL,
  PRIMARY KEY tag (netlocid,urlid,name(100),value(100))
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE search (
  recordid int(11) NOT NULL default '0',
  stext mediumtext,
  PRIMARY KEY (recordid),
  FULLTEXT (stext)
) ENGINE=MyISAM AVG_ROW_LENGTH = 20000 MAX_ROWS = 10000000 DEFAULT CHARACTER SET=utf8;
```

A.4.4   Administrative tables

```
CREATE TABLE netlocalias (
  netlocid int(11),
  netlocstr varchar(150) NOT NULL,
  KEY netlocid (netlocid),
  PRIMARY KEY netlocstr (netlocstr)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE urlalias (
  urlid int(11),
  urlstr tinytext,
  KEY urlid (urlid),
  PRIMARY KEY urlstr (urlstr(255))
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;
```

topichierarchy have to initialized manually

```
CREATE TABLE topichierarchy (
  node int(11) NOT NULL DEFAULT '0',
  father int(11) DEFAULT NULL,
  notation varchar(50) NOT NULL DEFAULT '',
  caption varchar(255) DEFAULT NULL,
  level int(11) DEFAULT NULL,
  PRIMARY KEY node (node),
  KEY father (father),
  KEY notation (notation)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE netlocs (
  netlocid int(11) NOT NULL auto_increment,
  netlocstr varchar(150) NOT NULL,
  retries int(11) NOT NULL DEFAULT 0,
  PRIMARY KEY (netlocstr),
  UNIQUE INDEX netlockid (netlocid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE urls (
  netlocid int(11) NOT NULL DEFAULT '0',
  urlid int(11) NOT NULL auto_increment,
  urlstr tinytext,
  path tinytext,
  PRIMARY KEY urlstr (urlstr(255)),
  INDEX netlocid (netlocid),
  UNIQUE INDEX urlid (urlid)
) ENGINE=MyISAM MAX_ROWS = 1000000000 DEFAULT CHARACTER SET=utf8;

CREATE TABLE urldb (
  netlocid int(11) NOT NULL default '0',
  netloclock int(11) NOT NULL default '0',
  urlid int(11) NOT NULL default '0',
  urllock int(11) NOT NULL default '0',
  harvest tinyint(1) NOT NULL default '0',
  retries int(11) NOT NULL default '0',
  PRIMARY KEY  (urlid),
  KEY netlocid (netlocid),
  KEY harvest (harvest)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE newlinks (
  urlid int(11) NOT NULL,
  netlocid int(11) NOT NULL,
  PRIMARY KEY  (urlid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;
```

```
CREATE TABLE recordurl (
  recordid int(11) NOT NULL auto_increment,
  urlid int(11) NOT NULL default '0',
  lastchecked timestamp NOT NULL,
  md5 char(32),
  fingerprint char(50),
  KEY md5 (md5),
  KEY fingerprint (fingerprint),
  PRIMARY KEY (urlid),
  KEY recordid (recordid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE admin (
  status enum('closed','open','paused','stopped') default NULL,
  schedulealgorithm enum('default','bigdefault','advanced') default 'default',
  queid int(11) NOT NULL default '0'
) ENGINE=MEMORY DEFAULT CHARACTER SET=utf8;

    Initialise admin to 'open' status
INSERT INTO admin VALUES ('open','default',0)


CREATE TABLE log (
  pid int(11) NOT NULL default '0',
  id varchar(50) default NULL,
  date timestamp NOT NULL,
  message varchar(255) default NULL
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE que (
  netlocid int(11) NOT NULL default '0',
  urlid int(11) NOT NULL default '0',
  queid int(11) NOT NULL auto_increment,
  PRIMARY KEY  (queid)
) ENGINE=MEMORY DEFAULT CHARACTER SET=utf8;

CREATE TABLE robotrules (
  netlocid int(11) NOT NULL default '0',
  expire int(11) NOT NULL default '0',
  rule varchar(255) default '',
  KEY netlocid (netlocid)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;

CREATE TABLE oai (
  recordid int(11) NOT NULL default '0',
  md5 char(32),
  date timestamp,
  status enum('created', 'updated', 'deleted'),
  PRIMARY KEY (md5),
  KEY date (date)
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;
```

```
CREATE TABLE exports (
  host varchar(30),
  port int,
  last timestamp DEFAULT '1999-12-31'
) ENGINE=MyISAM DEFAULT CHARACTER SET=utf8;
```

A.4.5   Create user dbuser with required priviligies

```
GRANT SELECT,INSERT,UPDATE,DELETE,CREATE,CREATE TEMPORARY TABLES,
    ALTER,LOCK TABLES ON $database.* TO $dbuser;

GRANT SELECT,INSERT,UPDATE,DELETE,CREATE,CREATE TEMPORARY TABLES,
    ALTER,LOCK TABLES ON $database.* TO $dbuser\@localhost;
```

A.5   Manual pages

A.5.1   combineExport

NAME   combineExport - export records in XML from Combine database

SYNOPSIS   combineExport –jobname <name> [–profile alvis|dc|combine –charset utf8|isolatin –number <n> –recordid <n> –md5 <MD5> –pipehost <server> –pipeport <n> –incremental]

OPTIONS AND ARGUMENTS   jobname is used to find the appropriate configuration (mandatory)

–profile

> Three profiles: alvis, dc, and combine . alvis and combine are similar XML formats.

> 'alvis' profile format is defined by the Alvis enriched document format DTD. It uses charset UTF-8 per default.

> 'combine' is more compact with less redundancy.

> 'dc' is XML encoded Dublin Core data.

–charset

> Selects a specific characterset from UTF-8, iso-latin-1 Overrides –profile settings.

–pipehost, –pipeport

> Specifies the server-name and port to connect to and export data using the Alvis Pipeline. Exports incrementally, ie all changes since last call to combineExport with the same pipehost and pipeport.

–number

> the max number of records to be exported

–recordid

> Export just the one record with this recordid

–md5

> Export just the one record with this MD5 checksum

–incremental

> Exports incrementally, ie all changes since last call to combineExport using –incremental

–xsltscript

> Generates records in Combine native format and converts them using this XSLT
> script before output. See example scripts in /etc/combine/*.xsl

DESCRIPTION

EXAMPLES

```
Export all records in Alvis XML-format to the file recs.xml
  combineExport --jobname atest > recs.xml

Export 10 records to STDOUT
  combineExport --jobname atest --number 10

Export all records in UTF-8 using Combine native format
  combineExport --jobname atest --profile combine --charset utf8 > Zebrarecs.xml

Incremental export of all changes from last call using localhost at port 6234 using the
default profile (Alvis)
  combineExport --jobname atest --pipehost localhost --pipeport 6234
```

SEE ALSO   Combine configuration documentation in */usr/share/doc/combine/*.
Alvis XML schema (–profile alvis) at http://project.alvis.info/alvis_docs/enriched-document.xsd

AUTHOR   Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2005 - 2006 Anders Ardö
This library is free software; you can redistribute it and/or modify it under the same
terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5
you may have available.

```
See the file LICENCE included in the distribution at
L<http://combine.it.lth.se/>
```

A.5.2   combineCtrl

NAME   combineCtrl - controls a Combine crawling job

SYNOPSIS   combineCtrl <action> –jobname <name>
where action can be one of start, kill, load, recyclelinks, reharvest, stat, howmany,
records, hosts, initMemoryTables, open, stop, pause, continue

OPTIONS AND ARGUMENTS   jobname is used to find the appropriate configuration (mandatory)

Actions starting/killing crawlers

start

takes an optional switch `-harvesters n` where `n` is the number of crawler processes to start

kill

kills all active crawlers (and their associated combineRun monitors) for jobname

Actions loading or recycling URLs for crawling

load

Read a list of URLs from STDIN (one per line) and schedules them for crawling

recyclelinks

Schedule all newly found (since last invocation of recyclelinks) links in crawled pages for crawling

reharvest

Schedules all pages in the database for crawling again (in order to check if they have changed)

Actions for controlling scheduling of URLs

open

opens database for URL scheduling (maybe after a stop)

stop

stops URL scheduling

pause

pauses URL scheduling

continue

continues URL scheduling after a pause

Misc actions

stat

prints out rudimentary status of the ready queue (ie eligible now) of URLs to be crawled

howmany

prints out rudimentary status of all URLs to be crawled

records

>   prints out the number of ercords in the SQL database

hosts

>   prints out rudimentary status of all hosts that have URLs to be crawled

initMemoryTables

>   initializes the administrative MySQL tables that are kept in memory

DESCRIPTION   Implements various control functionality to administer a crawling job, like starting and stoping crawlers, injecting URLs into the crawl queue, scheduling newly found links for crawling, controlling scheduling, etc.

This is the preferred way of controling a crawl job.

EXAMPLES

`echo ’http://www.yourdomain.com/’ | combineCtrl load -jobname aatest`

>   Seed the crawling job `aatest` with a URL

`combineCtrl start -jobname aatest -harvesters 3`

>   Start 3 crawling processes for job `aatest`

`combineCtrl recyclelinks -jobname aatest`

>   Schedule all new links crawling

`combineCtrl stat -jobname aatest`

>   See how many URLs that are eligible for crawling right now.

SEE ALSO   combine
    Combine configuration documentation in */usr/share/doc/combine/*.

AUTHOR   Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2005 Anders Ardö
    This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.
    See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.3   combineRun

NAME   combineRun - starts, monitors and restarts a combine harvesting process

SYNOPSIS   combineRun <pidfile> <combine command to run>

DESCRIPTION Starts a program and monitors it in order to make sure there is always a copy running. If the program dies it will be restarted with the same parameters. Used by `combineCtrl` when starting combine crawling.

SEE ALSO combineCtrl

AUTHOR Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE Copyright (C) 2005 Anders Ardö
    This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.
    See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.4 combineReClassify

NAME combineReClassify - main program that reanalyse records in a combine database
    Algorithm: select relevant records based on cls parameter for each record get record from database delete analyse infor from the record analyse the record if still_relevant save in database

---

A.5.5 combineSVM

NAME combineSVM - generate a SVM model from good and bad examples

SYNOPSIS combineSVM –jobname <name> [–good <good-file>] [–bad <bad-file>] [–train <model-file>] [–help]

OPTIONS AND ARGUMENTS jobname is used to find the appropriate configuration (mandatory)
    good is the name of a file with good URLs, one per line. Default 'goodURL.txt'
    bad is the name of a file with bad URLs, one per line. Default 'badURL.txt'
    train is the name of the file where the trained SVM model will be stored. Default 'SVMmodel.txt'

DESCRIPTION Takes two files, one with positive examples (good) and one with negative examples (bad) and trains a SVM classifier using these. The resulting model is stored in the file <train>.
    The example files should contain one URL per line and nothing else.

SEE ALSO combine
    Combine configuration documentation in *usr/share/doc/combine/*.

AUTHOR Ignacio Garcia Dorado Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2008 Ignacio Garcia Dorado, Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.6   combineRank

NAME   combineRank - calculates various Ranks for a Combine crawled database

SYNOPSIS   combineRank <action> –jobname <name> –verbose

where action can be one of PageRank, PageRankBL, NetLocRank, and exportLinkGraph. Results on STDOUT.

OPTIONS AND ARGUMENTS   jobname is used to find the appropriate configuration (mandatory)

verbose enables printing of ranks to STDOUT as SQL INSERT statements

Actions calculating variants of PageRank

PageRank

calculate standard PageRank

PageRankBL

calculate PageRanks with backlinks added for each link

NetLocRank

calculate SiteRank for each site and a local DocRank for documents within each site. Global ranks are then calulated as SiteRank * DocRank

Actions exporting link data

exportLinkGraph

export linkgraph from Combine database

DESCRIPTION   Implements calculation of different variants of PageRank.

Results are written to STDOUT and can be huge for large databases.

Linkgraph is exported in ASCII as a sparse matrix, one row per line. First integer is the ID (urlid) of a page with links. The rest of integers on the line are IDs for pages linked to. Ie 121 5624 23416 51423 267178 means that page 121 links to pages 5624 23416 51423 267178

EXAMPLES

`combineRank -jobname aatest -verbose PageRankBL`

> calculate PageRank with backlinks, result on STDOUT

`combineRank -jobname aatest -verbose exportLinkGraph`

> export the linkgraph to STDOUT

SEE ALSO   combine
Combine configuration documentation in */usr/share/doc/combine/*.

AUTHOR   Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2006 Anders Ardö
This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at `http://combine.it.lth.se/`

A.5.7   combineUtil

NAME   combineUtil - various operations on the Combine database

SYNOPSIS   combineUtil <action> –jobname <name>
where action can be one of stats, termstat, classtat, sanity, all, serveralias, resetOAI, restoreSanity, deleteNetLoc, deletePath, deleteMD5, deleteRecordid, addAlias

OPTIONS AND ARGUMENTS   jobname is used to find the appropriate configuration (mandatory)

Actions listing statistics

stats

> Global statistics about the database

termstat

> generates statistics about the terms from topic ontology matched in documents (can be long output)

classtat

> generates statistics about the topic classes assigned to documents

Actions for sanity controlls

sanity

> Performs various sanity checks on the database

restoreSanity

> Deletes records which sanity checks finds insane

resetOAI

> Removes all history (ie 'deleted' records) from the OAI table. This is done by removing the OAI table and recreating it from the existing database.

Action all   Does the actions: stats, sanity, classtat, termstat

Actions for deleting records

deleteNetLoc

> Deletes all records matching the ','-separated list of server net-locations (server-names optionally with port) in the switch –netlocstr. Net-locations can include SQL wild cards ('%').

deletePath

> Deletes all records matching the ','-separated list of URl paths (excluding net-locations) in the switch –pathsubstr. Paths can include SQL wild cards ('%').

deleteMD5

> Delete the record which has the MD5 in switch –md5

deleteRecordid

> Delete the record which has the recordid in switch –recordid

Actions for handling server aliases

serverAlias

> Detect server aliases in the current database and do a 'addAlias' on each detected alias.

addAlias

> Manually add a serveralias to the system. Requires switches –aliases and –preferred

DESCRIPTION   Does various statistics generation as well as performing sanity checks on the database

EXAMPLES

`combineUtil termstat -jobname aatest`

> Generate matched term statistics

SEE ALSO combine

Combine configuration documentation in */usr/share/doc/combine/*.

AUTHOR Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE Copyright (C) 2005 Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.8 combine

NAME Combine - Focused Web crawler framework

SYNOPSIS combine –jobname <name> –logname <id>

OPTIONS AND ARGUMENTS jobname is used to find the appropriate configuration (mandatory)

logname is used as identifier in the log (in MySQL table log)

DESCRIPTION Does crawling, parsing, optional topic-check and stores in MySQL database Normally started with the `combineCtrl` command. Briefly it get's an URL from the MySQL database, which acts as a common coordinator for a Combine job. The Web-page is fetched, provided it passes the robot exclusion protocoll. The HTML ic cleaned using `Tidy` and parsed into metadata, headings, text, links and link achors. Then it is stored (optionaly provided a topic-check is passed to keep the crawler focused) in the MySQL database in a structured form.

A simple workflow for a trivial crawl job might look like:

```
  Initialize database and configuration
combineINIT --jobname aatest
  Enter some seed URLs from a file with a list of URLs
combineCtrl  load --jobname aatest < seedURLs.txt
  Start 2 crawl processes
combineCtrl  start --jobname aatest --harvesters 2

  For some time occasionally schedule new links for crawling
combineCtrl recyclelinks --jobname aatest
  or look at the size of the ready queue
combineCtrl stat --jobname aatest

  When satisfied kill the crawlers
combineCtrl kill --jobname aatest
  Export data records in a highly structured XML format
combineExport --jobname aatest
```

For more complex jobs you have to edit the job configuration file.

81

SEE ALSO   combineINIT, combineCtrl

Combine configuration documentation in */usr/share/doc/combine/*.

AUTHOR   Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2005 Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.9   Combine::PosMatcher

NAME   PosMatcher

DESCRIPTION   This a module in the DESIRE automatic classification system. Copyright 1999.

Exported routines: 1. Fetching text: These routines all extract texts from a document (either a Combine record, a Combine XWI datastructure or a WWW-page identified by a URL. They all return: $meta, $head, $text, $url, $title, $size $meta: Metadata from document $head: Important text from document $text: Plain text from document $url: URL of the document $title: HTML title of the document $size: The size of the document

```
   Common input parameters:
       $DoStem: 1=do stemming; 0=no stemming
       $stoplist: object pointer to a LoadTermList object with a stoplist loaded
       $simple: 1=do simple loading; 0=advanced loading (might induce errors)

 getTextXWI
     parameters: $xwi, $DoStem, $stoplist, $simple
       $xwi is a Combine XWI datastructure

 getTextURL
    parameters: $url, $DoStem, $stoplist, $simple
       $url is the URL for the page to extract text from
```

2. Term matcher accepts a text as a (reference) parameter, matches each term in Term against text Matches are recorded in an associative array with class as key and summed weight as value. Match parameters: $text, $termlist $text: text to match against the termlist $termlist: object pointer to a LoadTermList object with a termlist loaded output: %score: an associative array with classifications as keys and scores as values

3. Heuristics: sum scores down the classification tree to the leafs cleanEiTree parameters: %res - an associative array from Match output: %res - same array

AUTHOR   Anders Ardö, <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE Copyright (C) 2005,2006 Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.10 Combine::selurl

NAME selurl - Normalise and validate URIs for harvesting

INTRODUCTION Selurl selects and normalises URIs on basis of both general practice (hostname lowercasing, portnumber substsitution etc.) and Combine-specific handling (aplpying config_allow, config_exclude, config_serveralias and other relevant config settings).

The Config settings catered for currently are:

maxUrlLength - the maximum length of an unnormalised URL allow - Perl regular to identify allowed URLs exclude - Perl regular expressions to exclude URLs from harvesting serveralias - Aliases of server names sessionids - List sessionid markers to be removed

A selurl object can hold a single URL and has methods to obtain its subparts as defined in URI.pm, plus some methods to normalise and validate it in Combine context.

BUGS Currently, the only schemes supported are http, https and ftp. Others may or may not work correctly. For one thing, we assume the scheme has an internet hostname/port.

clone() will only return a copy of the real URI object, not a new selurl.

URI URI-escapes the strings fed into it by new() once. Existing percent signs in the input are left untouched, which implicates that:

(a) there is no risk of double-encoding; and

(b) if the original contained an inadvertent sequence that could be interpreted as an escape sequence, uri_unescape will not render the original input (e.g. url_with_%66_in_it goes whoop) If you know that the original has not yet been escaped and wish to safeguard potential percent signs, you'll have to escape them (and only them) once before you offer it to new().

A problem with URI is, that its object is not a hash we can piggyback our data on, so I had to resort to AUTOLOAD to emulate inheritance. I find this ugly, but well, this *is* Perl, so what'd you expect?

A.5.11 Combine::XWI

NAME XWI.pm - class for internal representation of a document record

SYNOPSIS

```
use Combine::XWI;
$xwi = new Combine::XWI;
```

```
#single value record variables
$xwi->server($server);

my $server = $xwi->server();

#original content
$xwi->content(\$html);

my $text = ${$xwi->content()};

#multiple value record variables
$xwi->meta_add($name1,$value1);
$xwi->meta_add($name2,$value2);

$xwi->meta_rewind;
my ($name,$content);
while (1) {
 ($name,$content) = $xwi->meta_get;
 last unless $name;
}
```

DESCRIPTION   Provides methods for storing and retrieving structured records representing crawled documents.

METHODS

   new()

   XXX($val)   Saves $val using AUTOLOAD. Can later be retrieved, eg

```
 $xwi->MyVar('My value');
 $t = $xwi->MyVar;
```

will set $t to 'My value'

   *_reset()   Forget all values.

   *_rewind()   *_get will start with the first value.

   *_add   stores values into the datastructure

   *_get   retrieves values from the datastructure

meta_reset() / meta_rewind() / meta_add() / meta_get()   Stores the content of Meta-tags

Takes/Returns 2 parameters: Name, Content

```
$xwi->meta_add($name1,$value1);
$xwi->meta_add($name2,$value2);

$xwi->meta_rewind;
my ($name,$content);
while (1) {
 ($name,$content) = $xwi->meta_get;
 last unless $name;
}
```

xmeta_reset() / xmeta_rewind() / xmeta_add() / xmeta_get()   Extended information from Meta-tags. Not used.

url_remove() / url_reset() / url_rewind() / url_add() / url_get()   Stores all URLs (ie if multiple URLs for the same page) for this record

Takes/Returns 1 parameter: URL

heading_reset() / heading_rewind() / heading_add() / heading_get()   Stores headings from HTML documents

Takes/Returns 1 parameter: Heading text

link_reset() / link_rewind() / link_add() / link_get()   Stores links from documents

Takes/Returns 5 parameters: URL, netlocid, urlid, Anchor text, Link type

robot_reset() / robot_rewind() / robot_add() / robot_get()   Stores calculated information, like genre, language, etc

Takes/Returns 2 parameters Name, Value. Both are strings with max length Name: 15, Value: 20

topic_reset() / topic_rewind() / topic_add() / topic_get()   Stores result of topic classification.

Takes/Returns 5 parameters: Class, Absolute score, Normalized score, Terms, Algorithm id

Class, Terms, and Algorithm id are strings with max lengths Class: 50, and Algorithm id: 25

Absolute score, and Normalized score are integers

Normalized score and Terms are optional and may be replaced with 0, and " respectively

SEE ALSO   Combine focused crawler main site http://combine.it.lth.se/

AUTHOR   Yong Cao <tsao@munin.ub2.lu.se> v0.05 1997-03-13

Anders Ardö, <anders.ardo@it.lth.se>

A.5.12   Combine::Matcher

NAME   Matcher

DESCRIPTION   This a module in the DESIRE automatic classification system. Copyright 1999. Modified in the ALVIS project. Copyright 2004

Exported routines: 1. Fetching text: These routines all extract texts from a document (either a Combine XWI datastructure or a WWW-page identified by a URL. They all return: $meta, $head, $text, $url, $title, $size $meta: Metadata from document $head: Important text from document $text: Plain text from document $url: URL of the document $title: HTML title of the document $size: The size of the document

```
Common input parameters:
     $DoStem: 1=do stemming; 0=no stemming
     $stoplist: object pointer to a LoadTermList object with a stoplist loaded
     $simple: 1=do simple loading; 0=advanced loading (might induce errors)

getTextXWI
   parameters: $xwi, $DoStem, $stoplist, $simple
     $xwi is a Combine XWI datastructure

getTextURL
  parameters: $url, $DoStem, $stoplist, $simple
     $url is the URL for the page to extract text from
```

2. Term matcher accepts a text as a (reference) parameter, matches each term in Term against text Matches are recorded in an associative array with class as key and summed weight as value. Match parameters: $text, $termlist $text: text to match against the termlist $termlist: object pointer to a LoadTermList object with a termlist loaded output: %score: an associative array with classifications as keys and scores as values

AUTHOR   Anders Ardö <anders.ardo@it.lth.se>

A.5.13   Combine::FromTeX

NAME   Combine::FromTeX.pm - TeX parser in combine package

AUTHOR

Anders Ardø 2000-06-11

A.5.14 Combine::SD_SQL

NAME SD_SQL

DESCRIPTION Reimplementation of sd.pl SD.pm and SDQ.pm using MySQL contains both recyc and guard

Basic idea is to have a table (urldb) that contains most URLs ever inserted into the system together with a lock (the guard function) and a boolean harvest-flag. Also in this table is the host part together with its lock. URLs are selected from this table based on urllock, netloclock and harvest and inserted into a queue (table que). URLs from this queue are then given out to harvesters. The queue is implemented as: # The admin table can be used to generate sequence numbers like this: #mysql> update admin set queid=LAST_INSERT_ID(queid+1); # and used to extract the next URL from the queue #mysql> select host,url from que where queid=LAST_INSERT_ID(); # When the queue is empty it is filled from table urldb. Several different algorithms can be used to fill it (round-robin, most urls, longest time since harvest, ...). Since the harvest-flag and guard-lock are not updated until the actual harvest is done it is OK to delete the queue and regenerate it anytime.

################################# #Questions, ideas, TODOs, etc #Split table urldb into 2 tables - one for urls and one for hosts??? #Less efficient when filling que; more efficient when updating netloclock #Datastruktur TABLE hosts: create table hosts(host varchar(50) not null default ", netloclock int not null, retries int not null default 0, ant int not null default 0, primary key (host), key (ant), key (netloclock));

############### Handle to many retries?

```
algorithm takes an url from the host that was accessed longest ago
($hostid,$url)=SELECT host,url,id FROM hosts,urls WHERE
     hosts.hostlock < UNIX_TIMESTAMP()
     hosts.host=urls.host AND
     urls.urllock < UNIX_TIMESTAMP() AND
     urls.harvest=1 ORDER BY hostlock LIMIT 1;

algorithm takes an url from the host with most URLs
($hostid,$url)=SELECT host,url,id FROM hosts,urls WHERE
     hosts.hostlock < UNIX_TIMESTAMP()
     hosts.host=urls.host AND
     urls.urllock < UNIX_TIMESTAMP() AND
     urls.harvest=1 ORDER BY host.ant DESC LIMIT 1;

algorithm takes an url from any available host
($hostid,$url)=SELECT host,url,id FROM hosts,urls WHERE
     hosts.hostlock < UNIX_TIMESTAMP()
     hosts.host=urls.host AND
     urls.urllock < UNIX_TIMESTAMP() AND
     urls.harvest=1 LIMIT 1;
```

AUTHOR   Anders Ardö <anders.ardo@it.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2005,2006 Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.15   Combine::utilPlugIn

NAME   utilPlugIn

DESCRIPTION   Utilities for: * extracting text from XWI's * SVM classification * language and country identification

AUTHOR   Ignacio Garcia Dorado Anders Ardö <anders.ardo@eit.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2008 Ignacio Garcia Dorado, Anders Ardö

This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.16   Combine::FromHTML

NAME   Combine::FromHTML.pm - HTML parser in combine package

AUTHOR   Yong Cao <tsao@munin.ub2.lu.se> v0.06 1997-03-19 Anders Ardø 1998-07-18 added <AREA ... HREF=link ...> fixed <A ... HREF=link ...> regexp to be more general Anders Ardö 2002-09-20 added 'a' as a tag not to be replaced with space added removal of Cntrl-chars and some punctuation marks from IP added <style>...</style> as something to be removed before processing beefed up compression of sequences of blanks to include \240 (non-breakable space) changed 'remove head' before text extraction to handle multiline matching (which can be introduced by decoding html entities) added compress blanks and remove CRs to metadata-content Anders Ardö 2004-04 Changed extraction process dramatically

---

A.5.17   Combine::RobotRules

NAME   RobotRules.pm

AUTHOR   Anders Ardo version 1.0 2004-02-19

A.5.18 Combine::HTMLExtractor

NAME HTMLExtractor

DESCRIPTION Adopted from HTML::LinkExtractor - Extract links from an HTML document by D.H (PodMaster)

AUTHOR Anders Ardo D.H (PodMaster)

LICENSE Copyright (c) 2003 by D.H. (PodMaster). All rights reserved.
  This module is free software; you can redistribute it and/or modify it under the same terms as Perl itself. The LICENSE file contains the full text of the license.

A.5.19 Combine::LoadTermList

NAME LoadTermList

DESCRIPTION This a module in the DESIRE automatic classification system. Copyright 1999.
  LoadTermList - A class for loading and storing a stoplist with single words a termlist with classifications and weights

```
 Subroutines:
   LoadStopWordList(StopWordListFileName)
      loads a list of stopwords, one per line, from
      the file StopWordListFileName.

   EraseStopWordList
      clears the stopword list

 Subroutines:
  LoadTermList(TermListFileName) - loads TermClass from file
  LoadTermListStemmed(TermListFileName) - same plus stems terms

 Input: A formatted term-list including weights and classifications
  Format:  <weight>: <term_reg_exp>=[<classification>, ]+
  weight can be a positive or negative number
  term_reg_exp can be words, phrases, boolean expressions (with @and
     as operator) on term_reg_exp or Perl regular expressions
```

AUTHOR Anders Ardö <Anders.Ardo@it.lth.se>

COPYRIGHT AND LICENSE Copyright (C) 2005,2006 Anders Ardö
  This library is free software; you can redistribute it and/or modify it under the same terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5 you may have available.
  See the file LICENCE included in the distribution at http://combine.it.lth.se/

A.5.20   Combine::classifySVM

NAME   classifySVM

DESCRIPTION   Classification plugin module using SVM (implementation SVMLight)
Uses SVM model loaded from file pointed to by configuration variable 'SVMmodel'

AUTHOR   Ignacio Garcia Dorado Anders Ardö <anders.ardo@eit.lth.se>

COPYRIGHT AND LICENSE   Copyright (C) 2008 Ignacio Garcia Dorado, Anders
Ardö

This library is free software; you can redistribute it and/or modify it under the same
terms as Perl itself, either Perl version 5.8.4 or, at your option, any later version of Perl 5
you may have available.

See the file LICENCE included in the distribution at http://combine.it.lth.se/